



THE STATE UNIVERSITY
OF NEW JERSEY

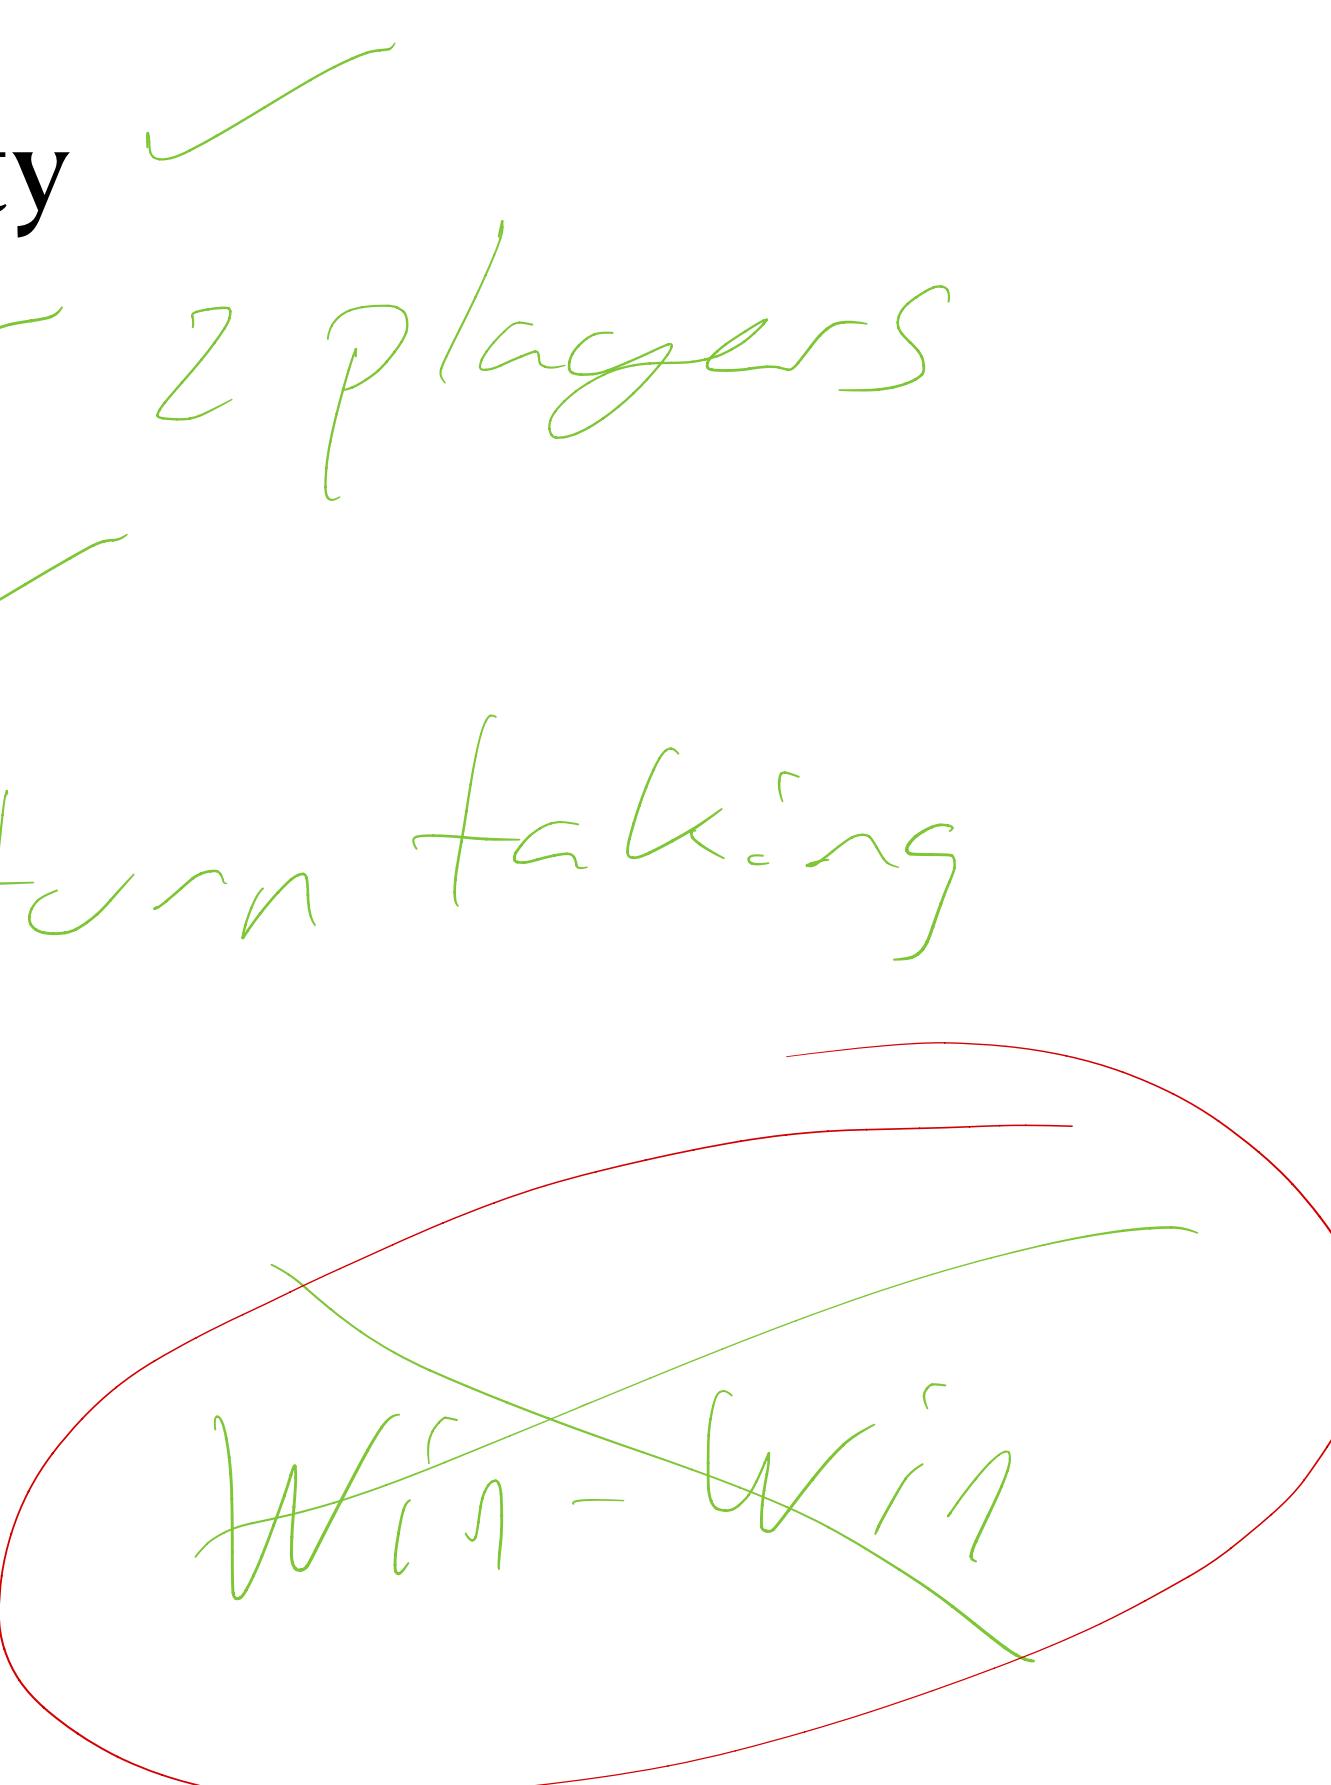
Adversarial Search Games

Edgar Granados

R

Games Properties

- Full Observability ✓
- Multi-agent ✓
- Deterministic ✓
- Sequential
- Discrete
- Known ✓
- Zero-sum

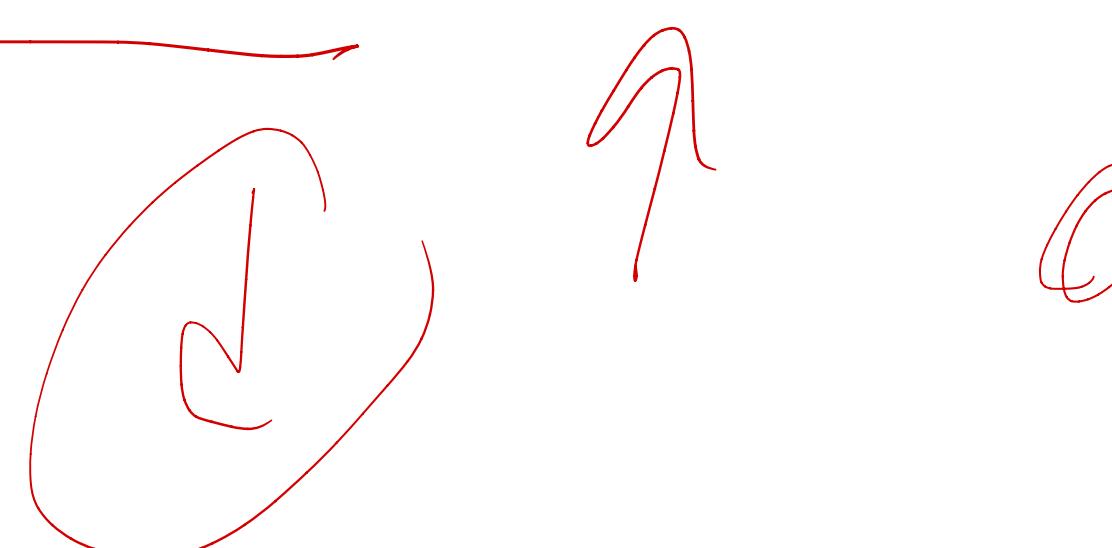


R

Game

- **Initial State (S_0)**: Where the game starts
- **To-Move(s)**: The player whose turn it is to move at state s
- **Actions**: Set of legal moves
- **Transition model**: The effect of applying the action
- **Is-Terminal(s)**: Returns true if the game is over
- **Utility(s,p)**: A function that assigns a numeric value to the final state

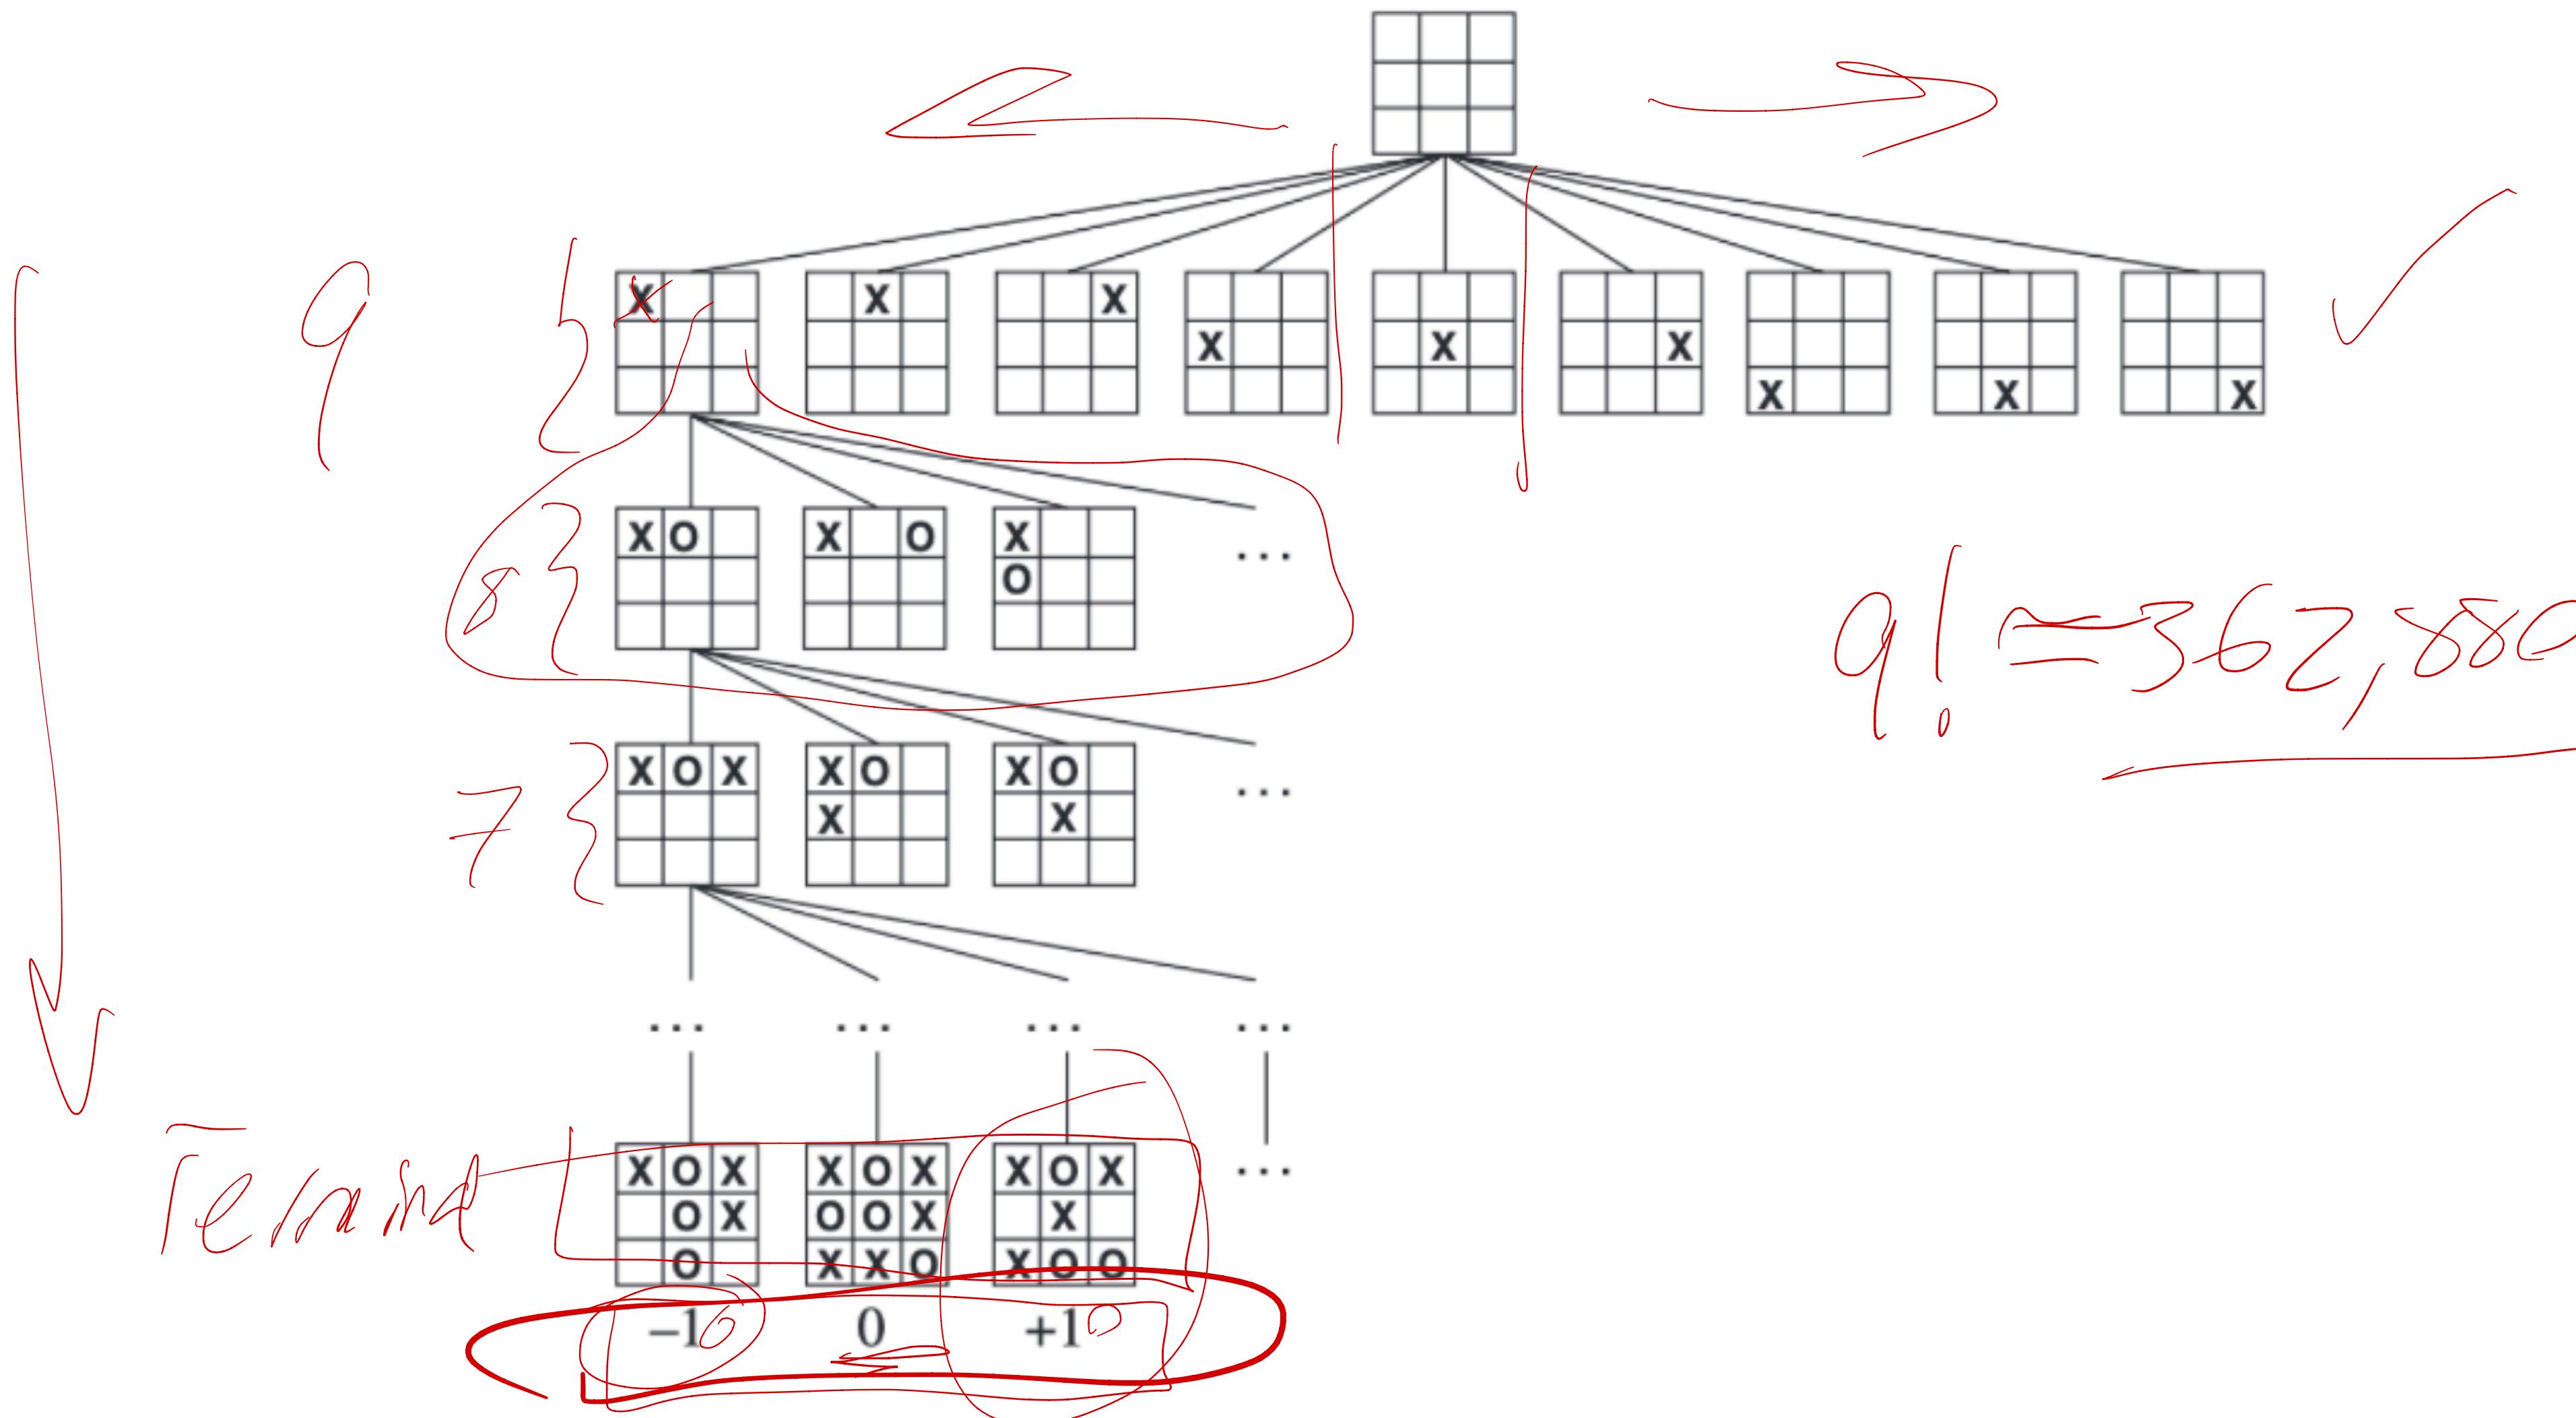
Deterministic



R

Examples

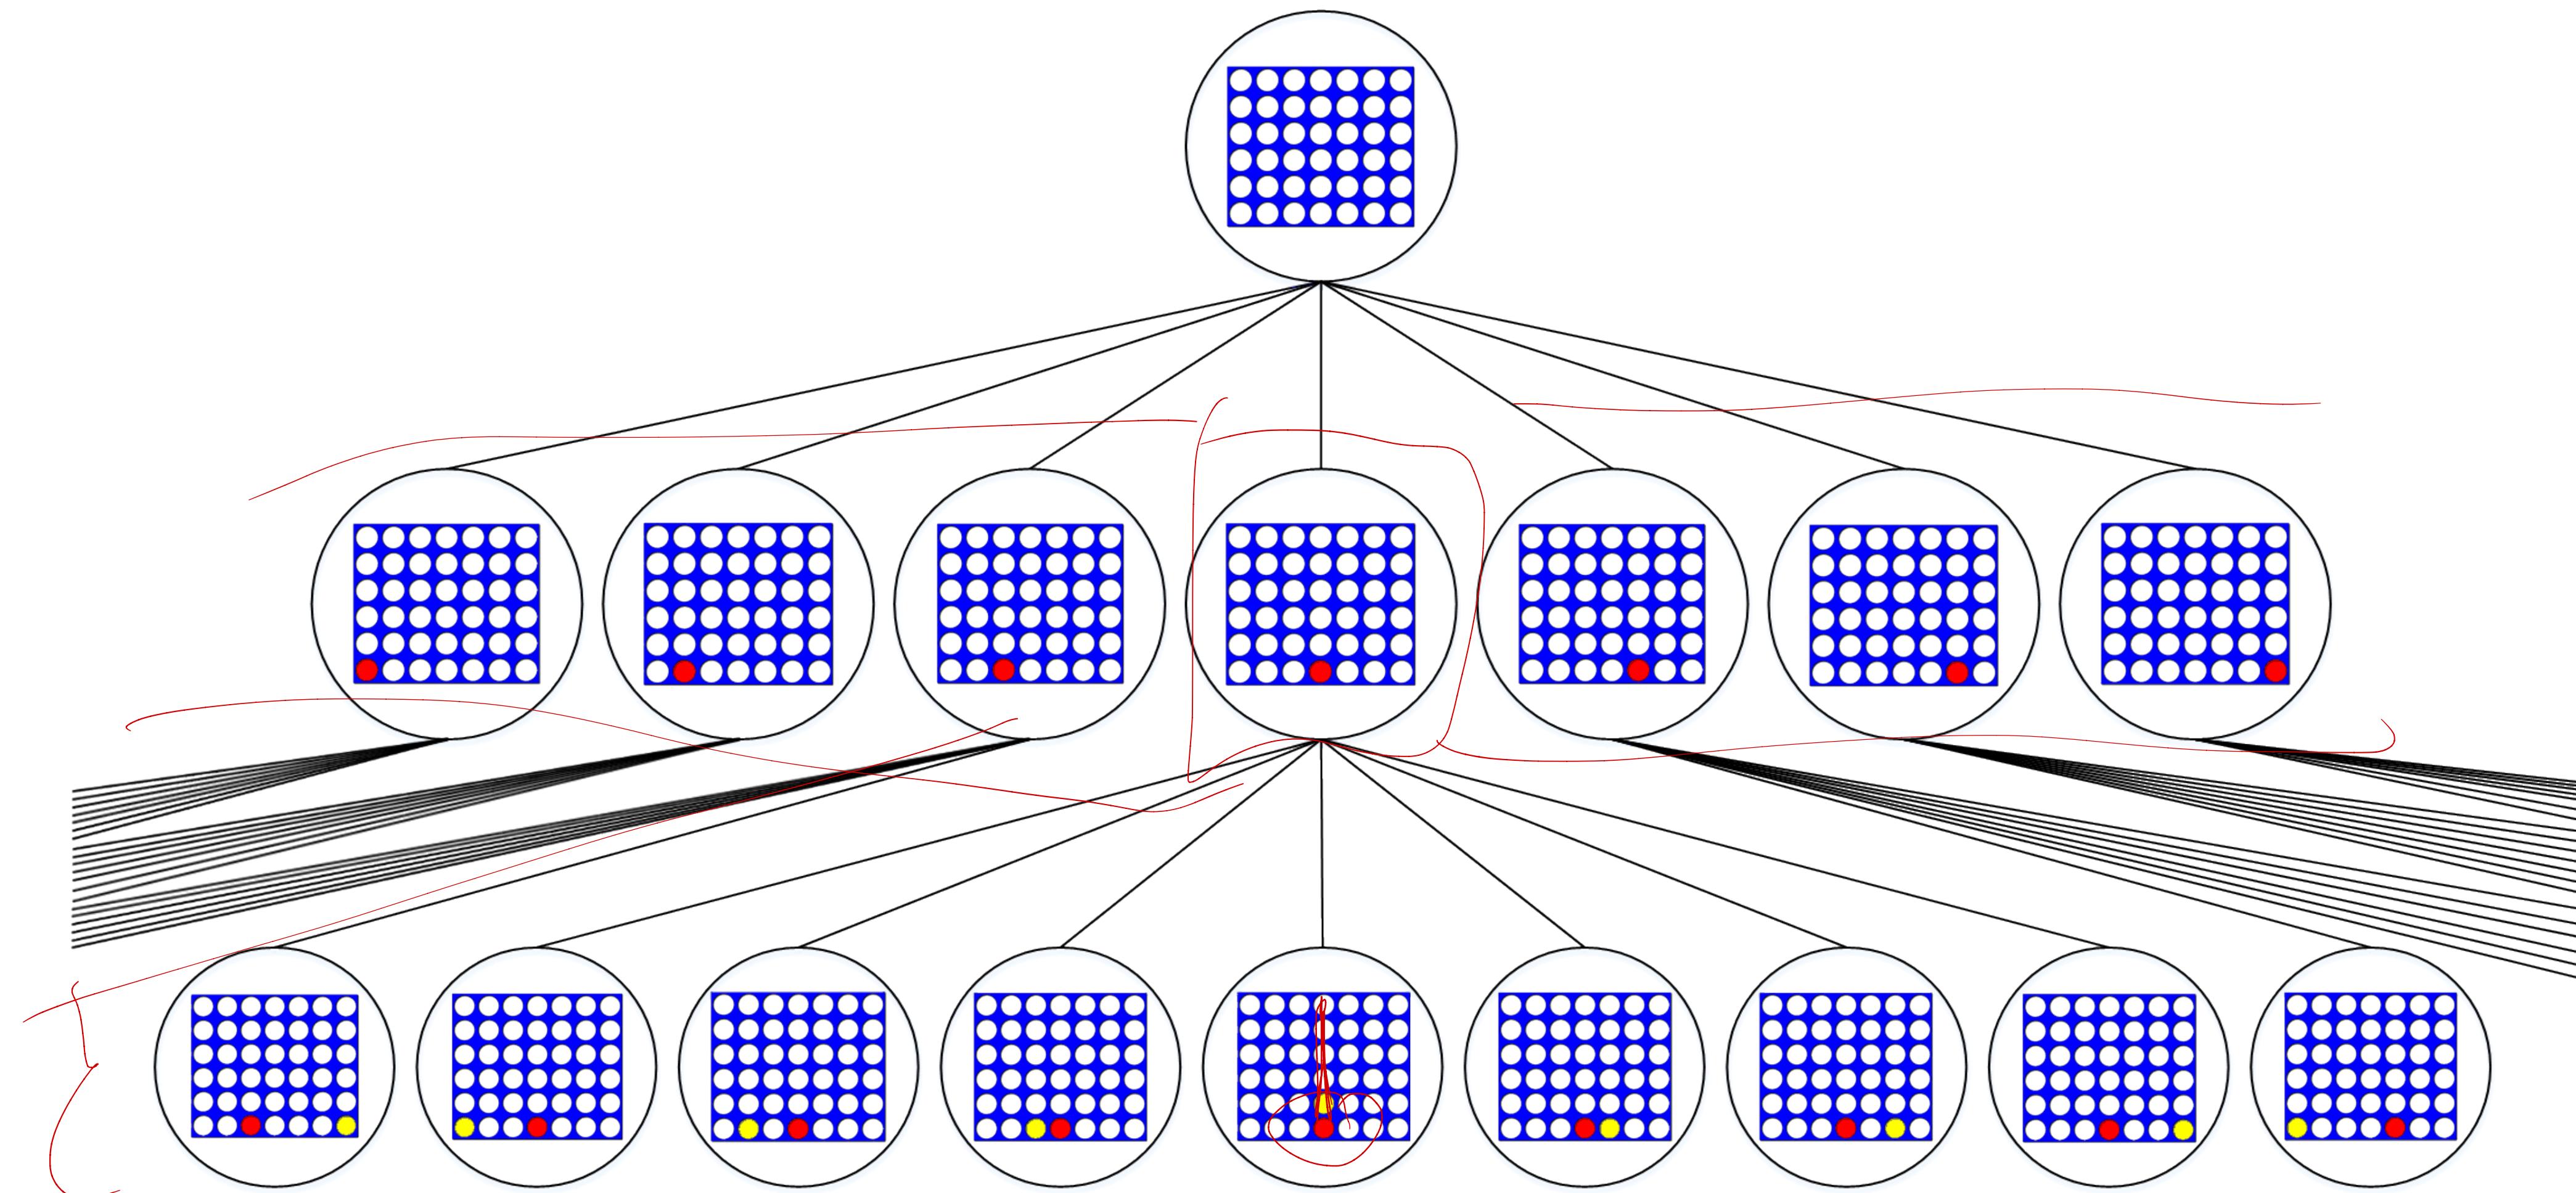
Tic-Tac-Toe



R

Examples

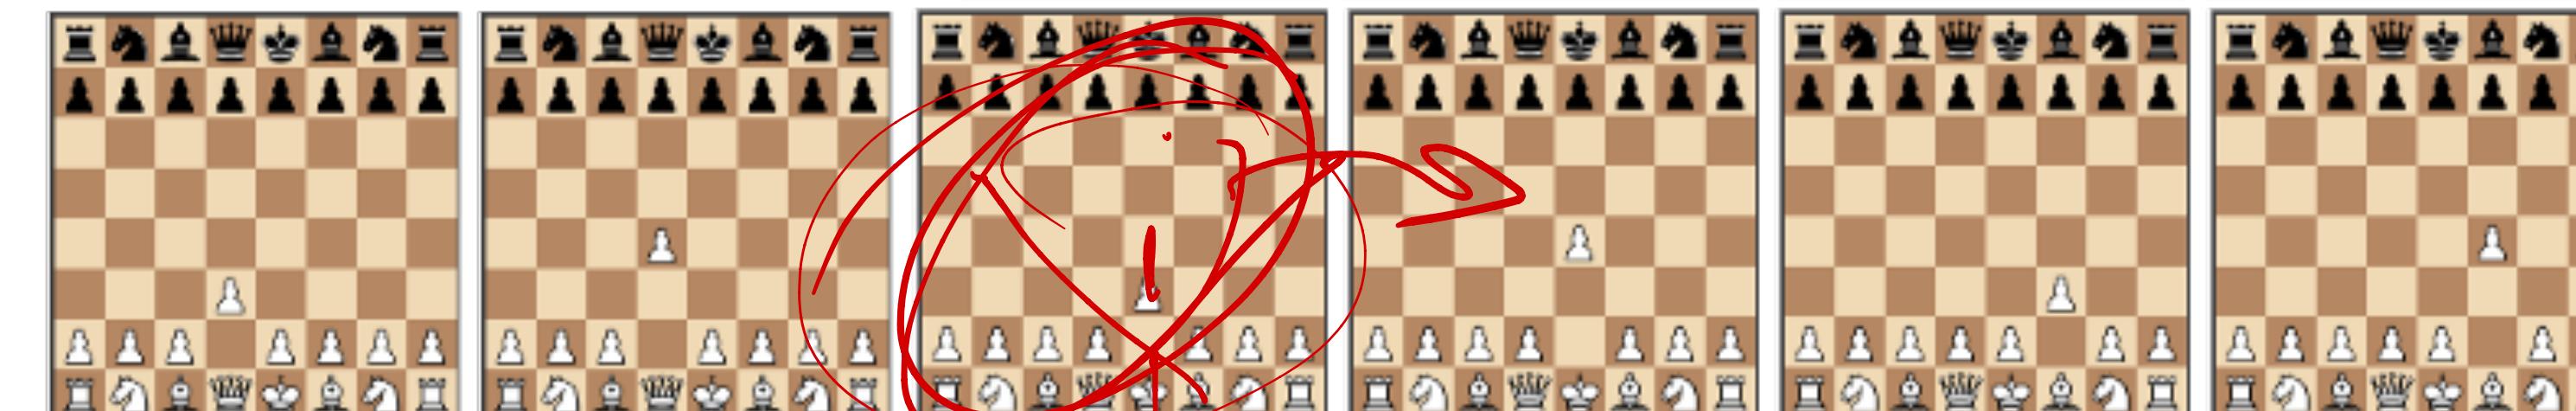
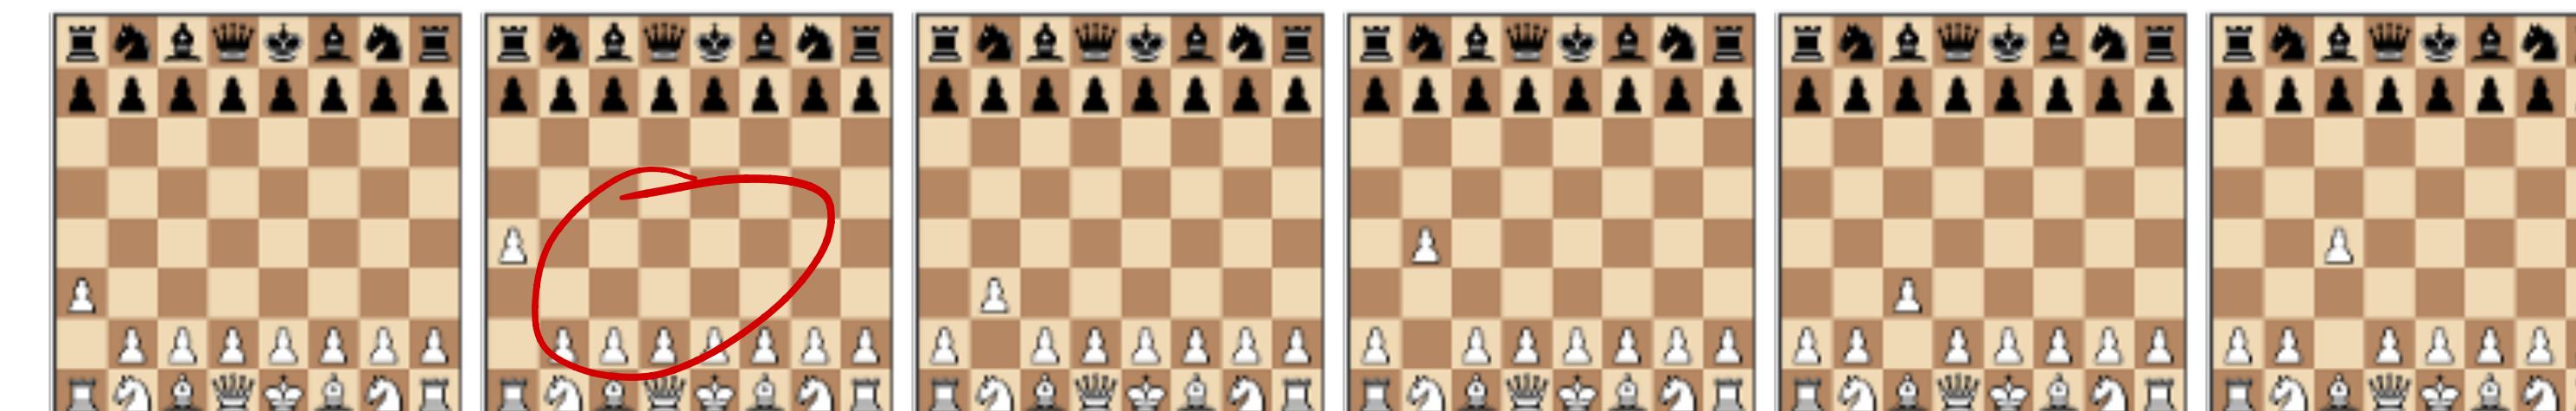
Connect 4



R

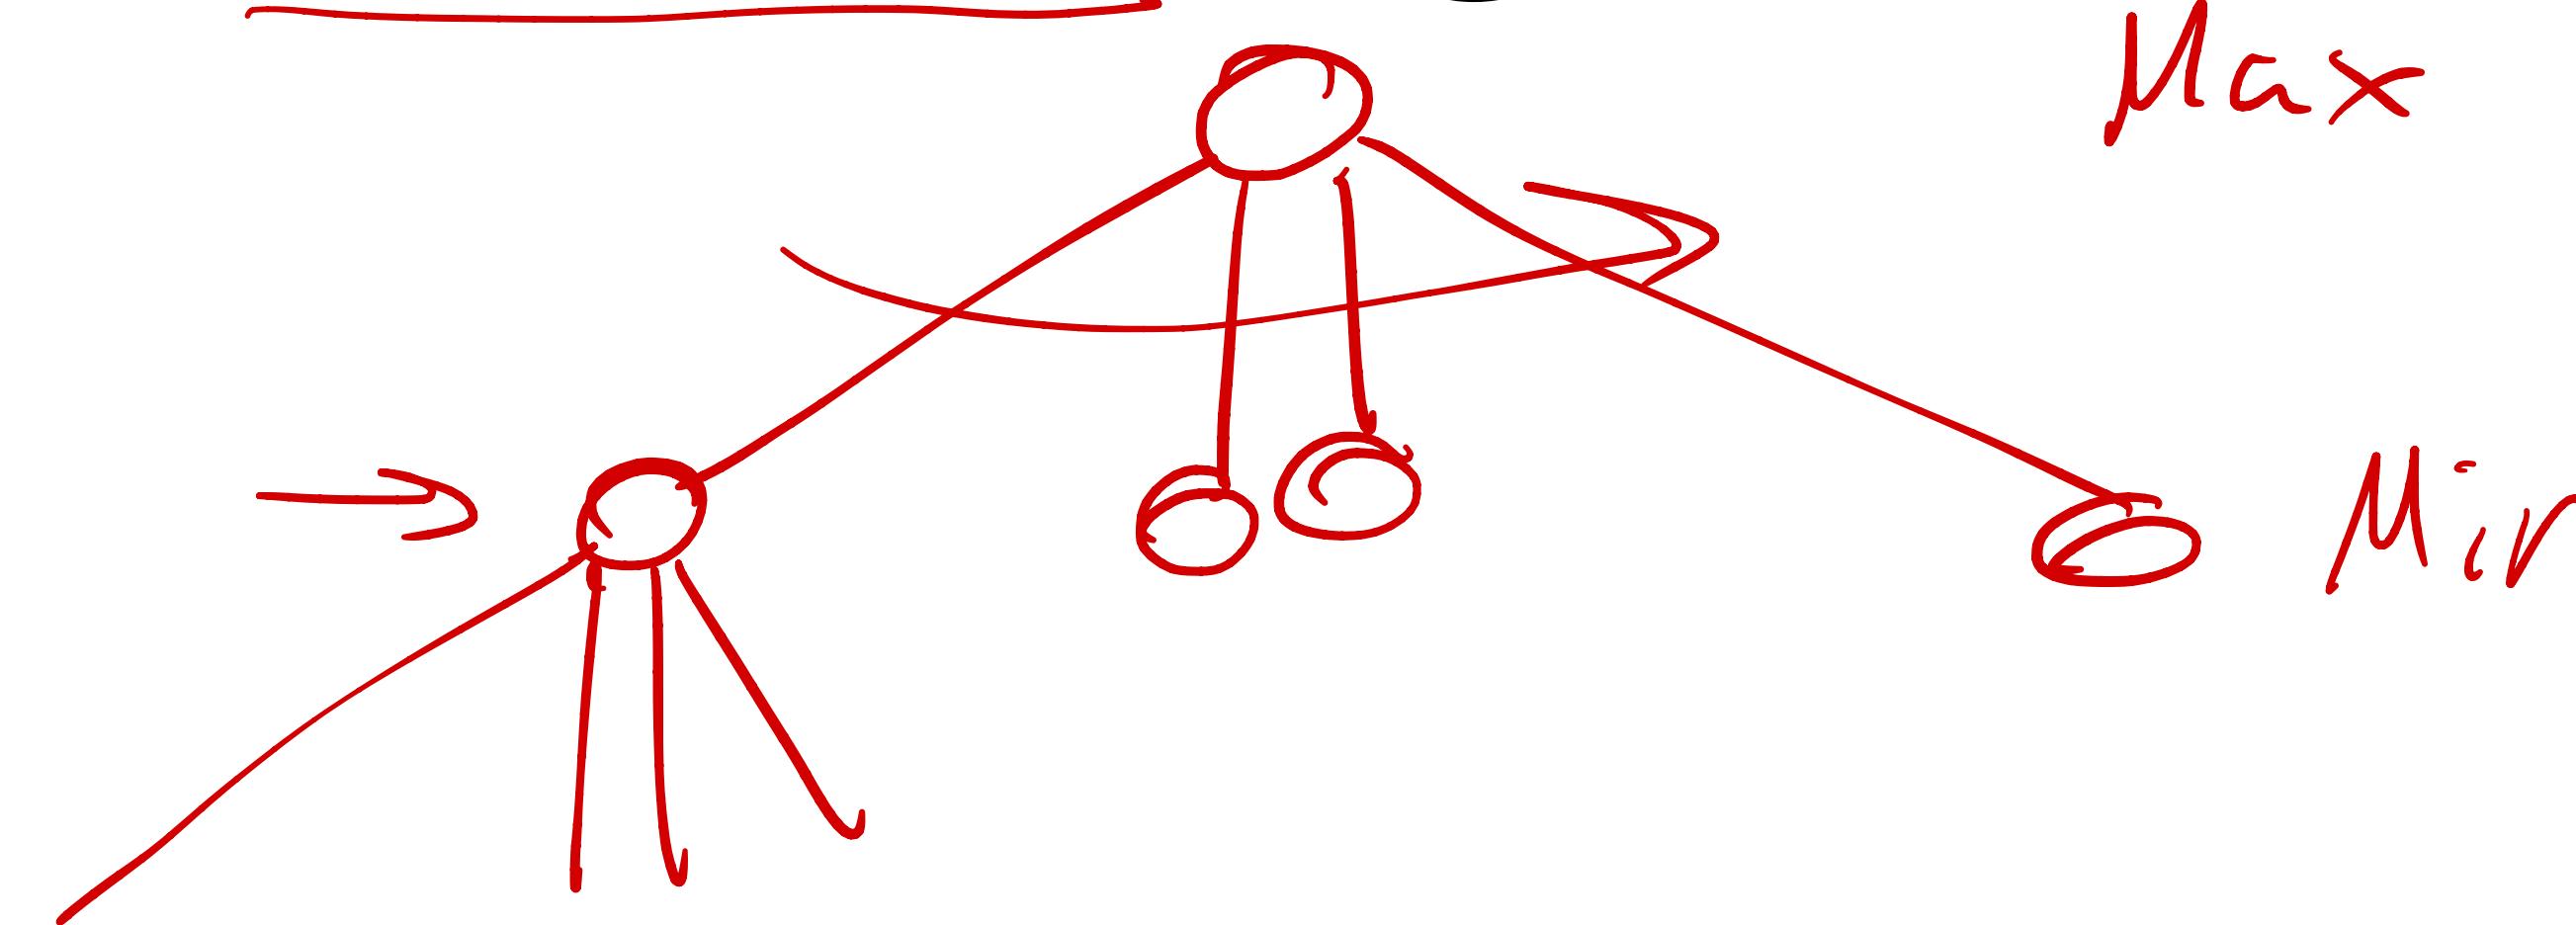
Examples

Chess



R

MinMax Algorithm



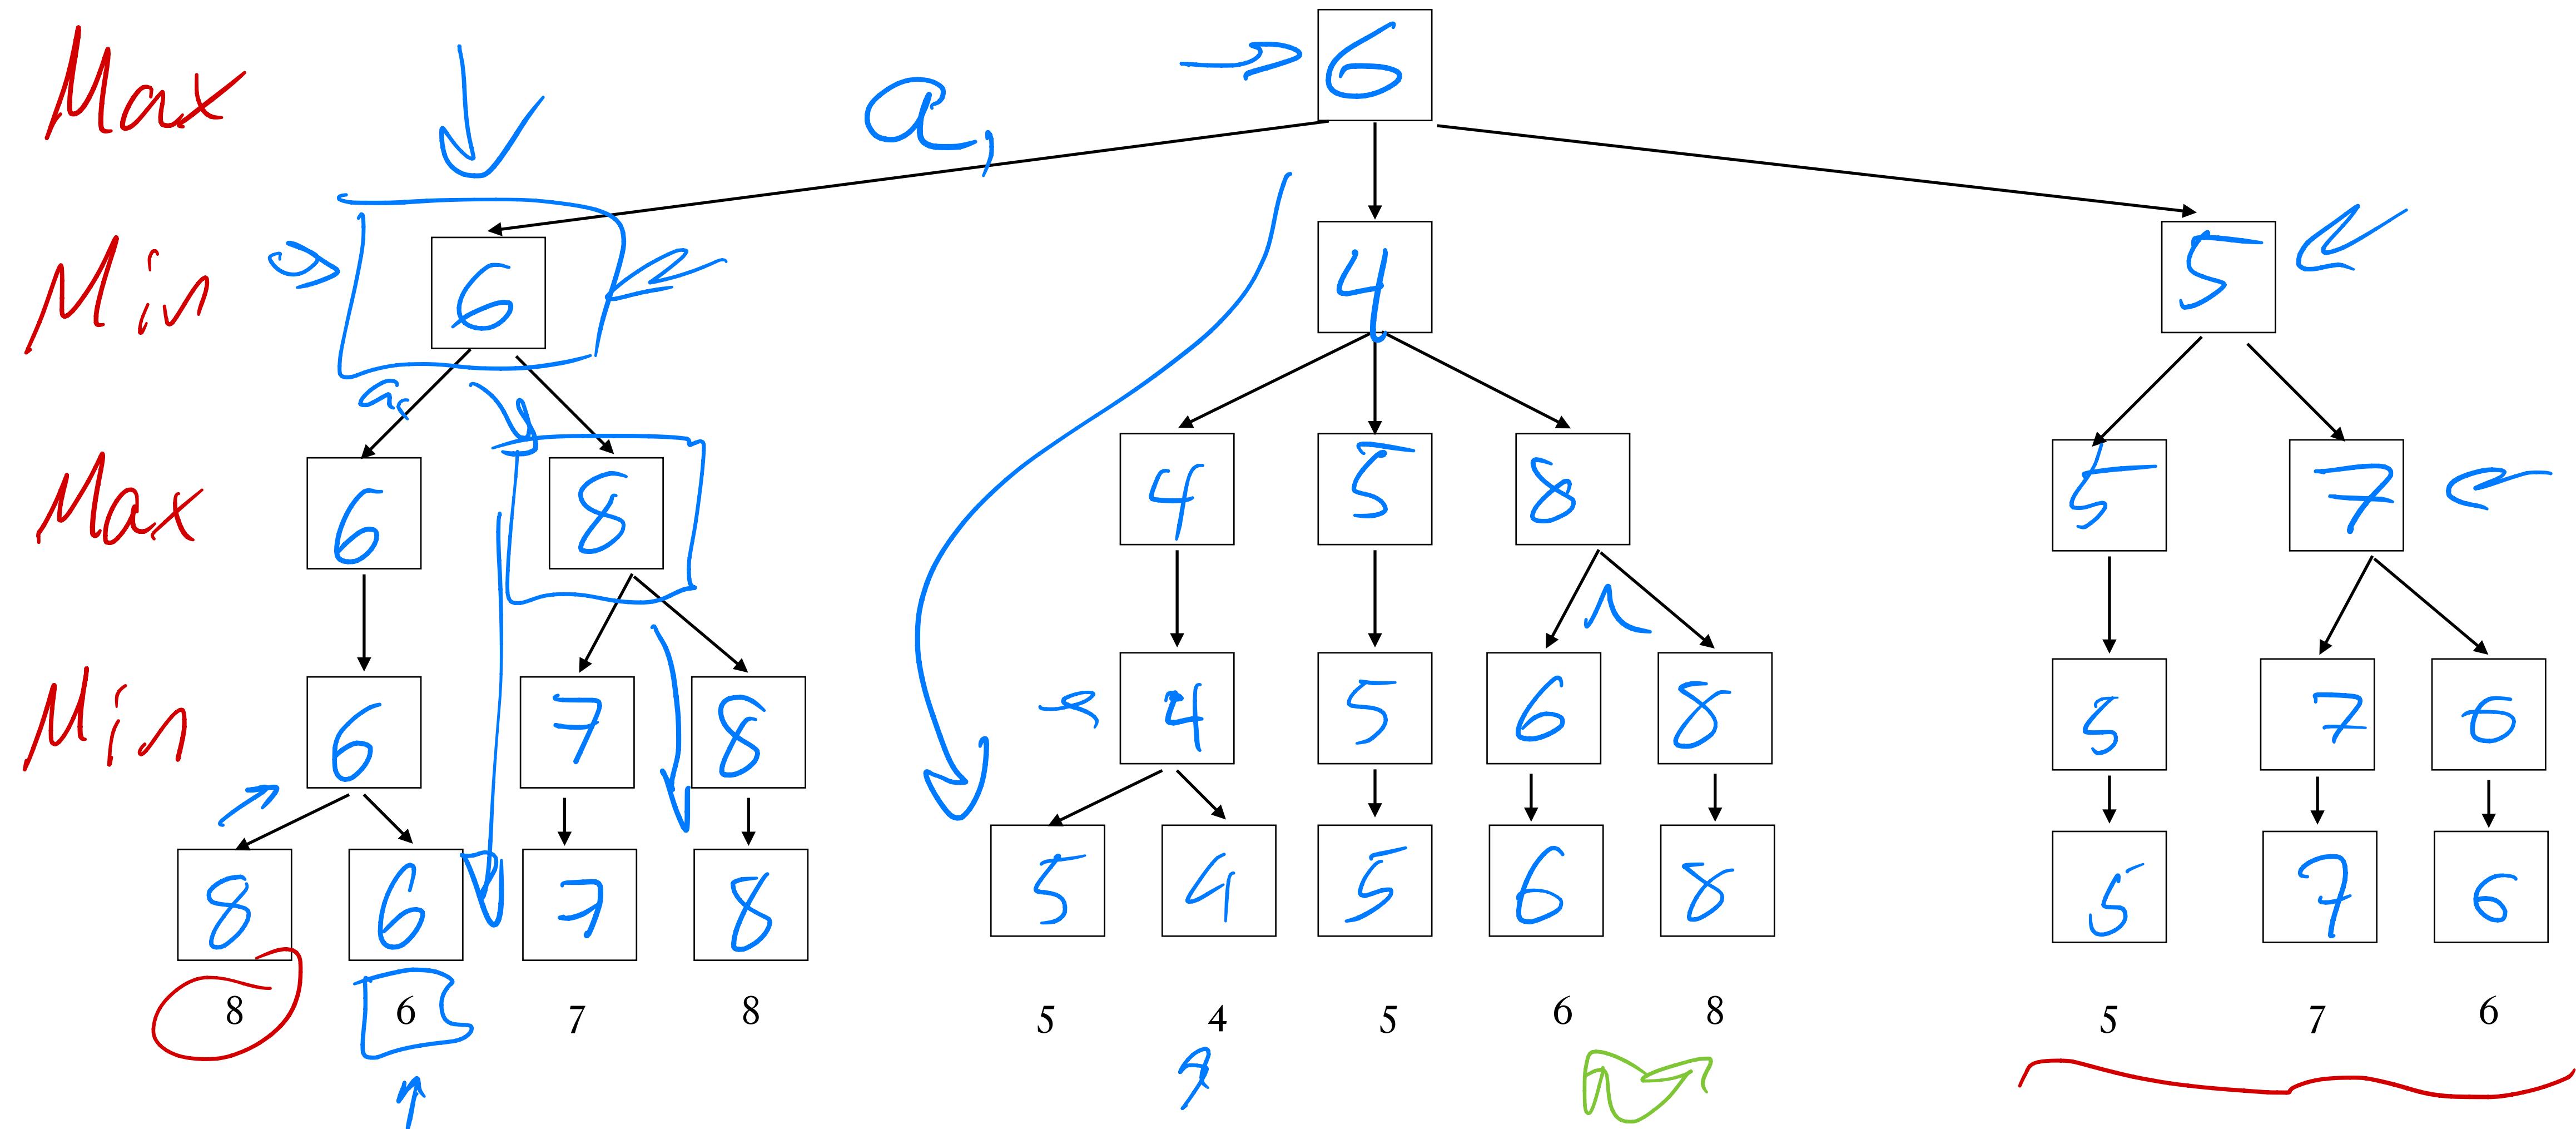
DFS

$$\text{MINIMAX}(s) = \begin{cases} \text{UTILITY}(s) & \text{if TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX}, \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN}. \end{cases}$$

R

MinMax Algorithm

Example: 2 players



R

MinMax Algorithm

Pseudocode

```
function MINIMAX-SEARCH(game, state) returns an action
    player  $\leftarrow$  game.TO-MOVE(state)
    value, move  $\leftarrow$  MAX-VALUE(game, state)
    return move

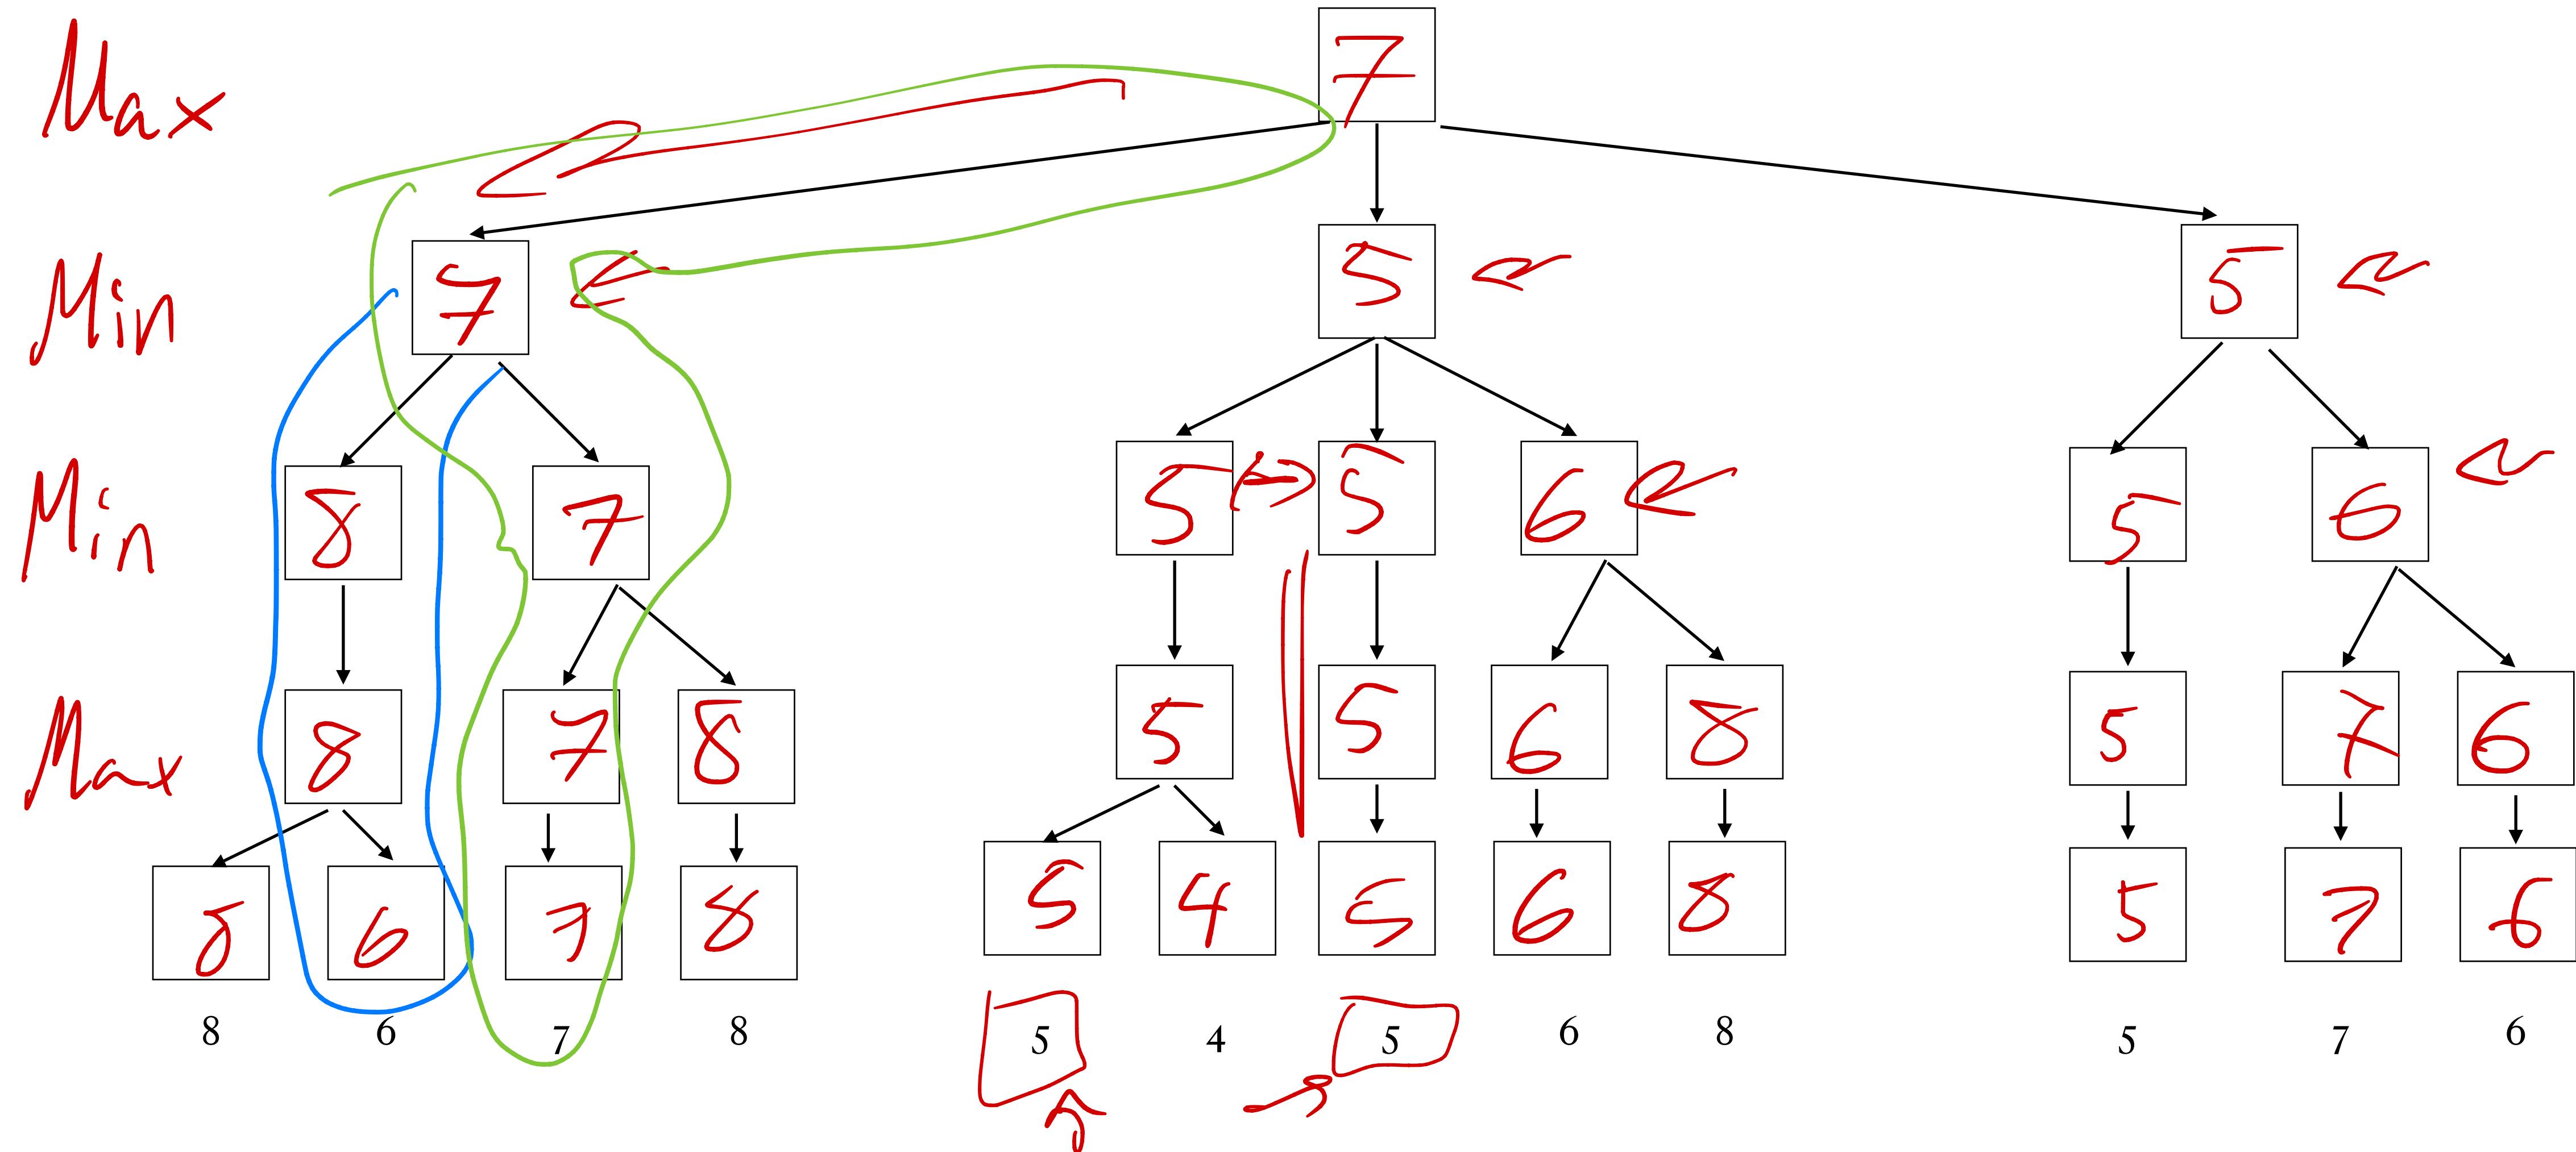
function MAX-VALUE(game, state) returns a (utility, move) pair
    if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
    v, move  $\leftarrow -\infty$ , null
    for each a in game.ACTIONS(state) do
        v2, a2  $\leftarrow$  MIN-VALUE(game, game.RESULT(state, a))
        if v2 > v then
            v, move  $\leftarrow$  v2, a
    return v, move

function MIN-VALUE(game, state) returns a (utility, move) pair
    if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
    v, move  $\leftarrow +\infty$ , null
    for each a in game.ACTIONS(state) do
        v2, a2  $\leftarrow$  MAX-VALUE(game, game.RESULT(state, a))
        if v2 < v then
            v, move  $\leftarrow$  v2, a
    return v, move
```

R

MinMax Algorithm

Example: 3 players



R

MinMax Algorithm

Alpha-Beta Pruning ↳ ~ 2 Players

Max

Min

Max

Min

8

6

7

1000

5

4

3

2

1

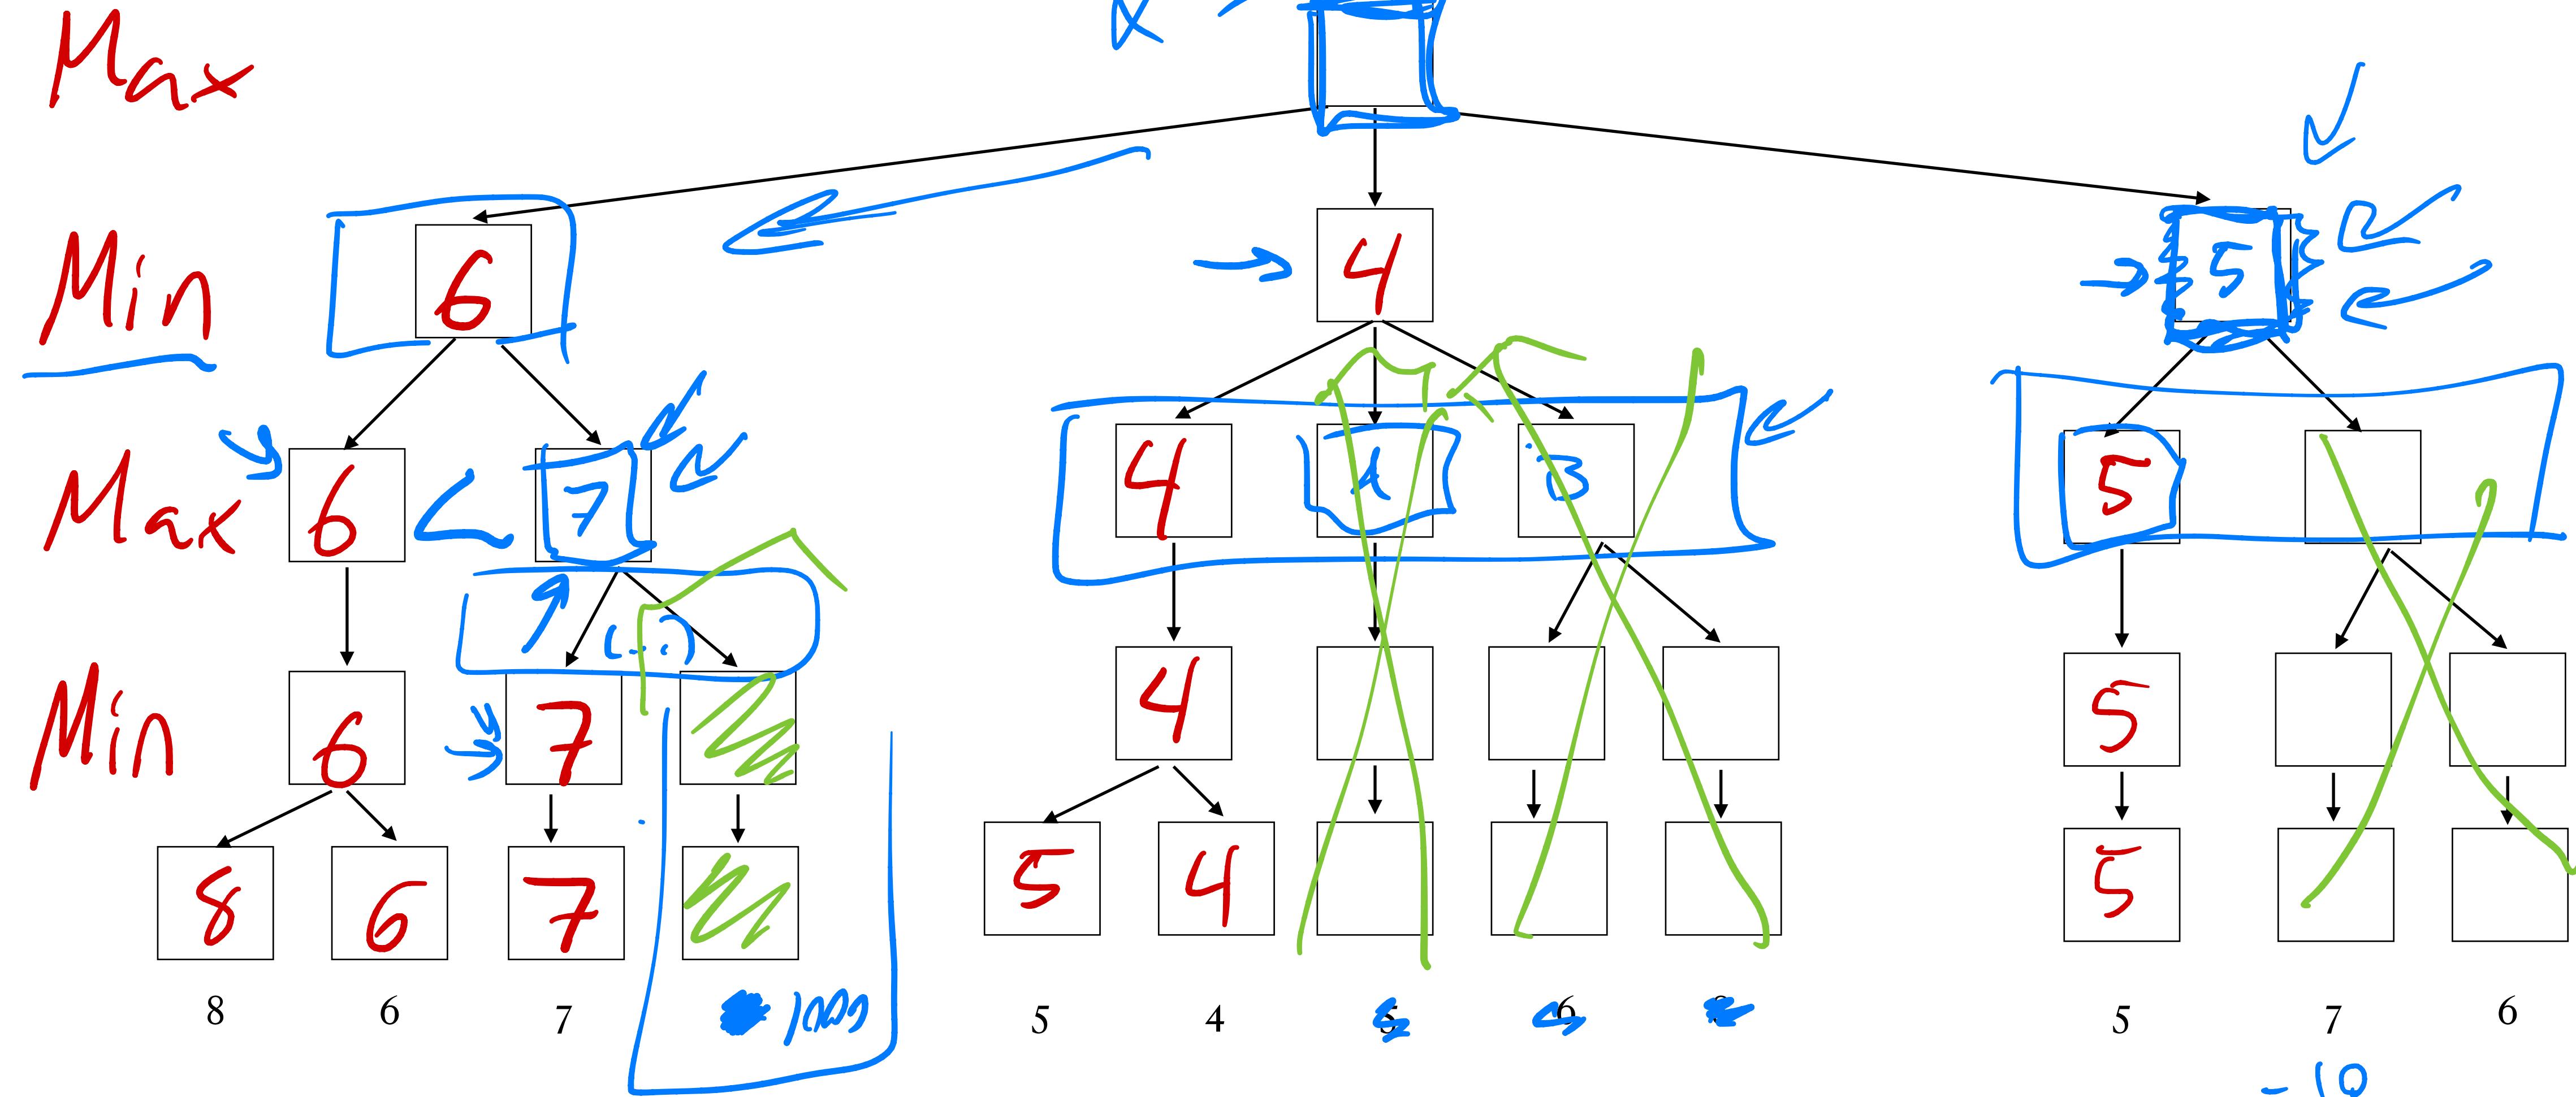
0

5

7

6

-10



R

MinMax Algorithm

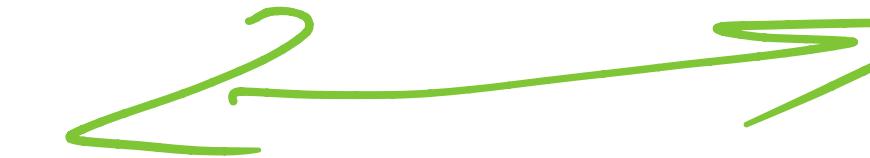
Alpha-Beta Pruning: Pseudocode

Original

```
function MINIMAX-SEARCH(game, state) returns an action
    player ← game.TO-MOVE(state)
    value, move ← MAX-VALUE(game, state)
    return move

function MAX-VALUE(game, state) returns a (utility, move) pair
    if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
    v, move ← -∞, null
    for each a in game.ACTIONS(state) do
        v2, a2 ← MIN-VALUE(game, game.RESULT(state, a))
        if v2 > v then
            v, move ← v2, a
    return v, move

function MIN-VALUE(game, state) returns a (utility, move) pair
    if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
    v, move ← +∞, null
    for each a in game.ACTIONS(state) do
        v2, a2 ← MAX-VALUE(game, game.RESULT(state, a))
        if v2 < v then
            v, move ← v2, a
    return v, move
```



α, β

```
function ALPHA-BETA-SEARCH(game, state) returns an action
    player ← game.TO-MOVE(state)
    value, move ← MAX-VALUE(game, state, -∞, +∞)
    return move

function MAX-VALUE(game, state, α, β) returns a (utility, move) pair
    if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
    v ← -∞
    for each a in game.ACTIONS(state) do
        v2, a2 ← MIN-VALUE(game, game.RESULT(state, a), α, β)
        if v2 > v then
            v, move ← v2, a
            α ← MAX(α, v)
    if v ≥ β then return v, move
    return v, move

function MIN-VALUE(game, state, α, β) returns a (utility, move) pair
    if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
    v ← +∞
    for each a in game.ACTIONS(state) do
        v2, a2 ← MAX-VALUE(game, game.RESULT(state, a), α, β)
        if v2 < v then
            v, move ← v2, a
            β ← MIN(β, v)
    if v ≤ α then return v, move
    return v, move
```

R

MinMax Algorithm

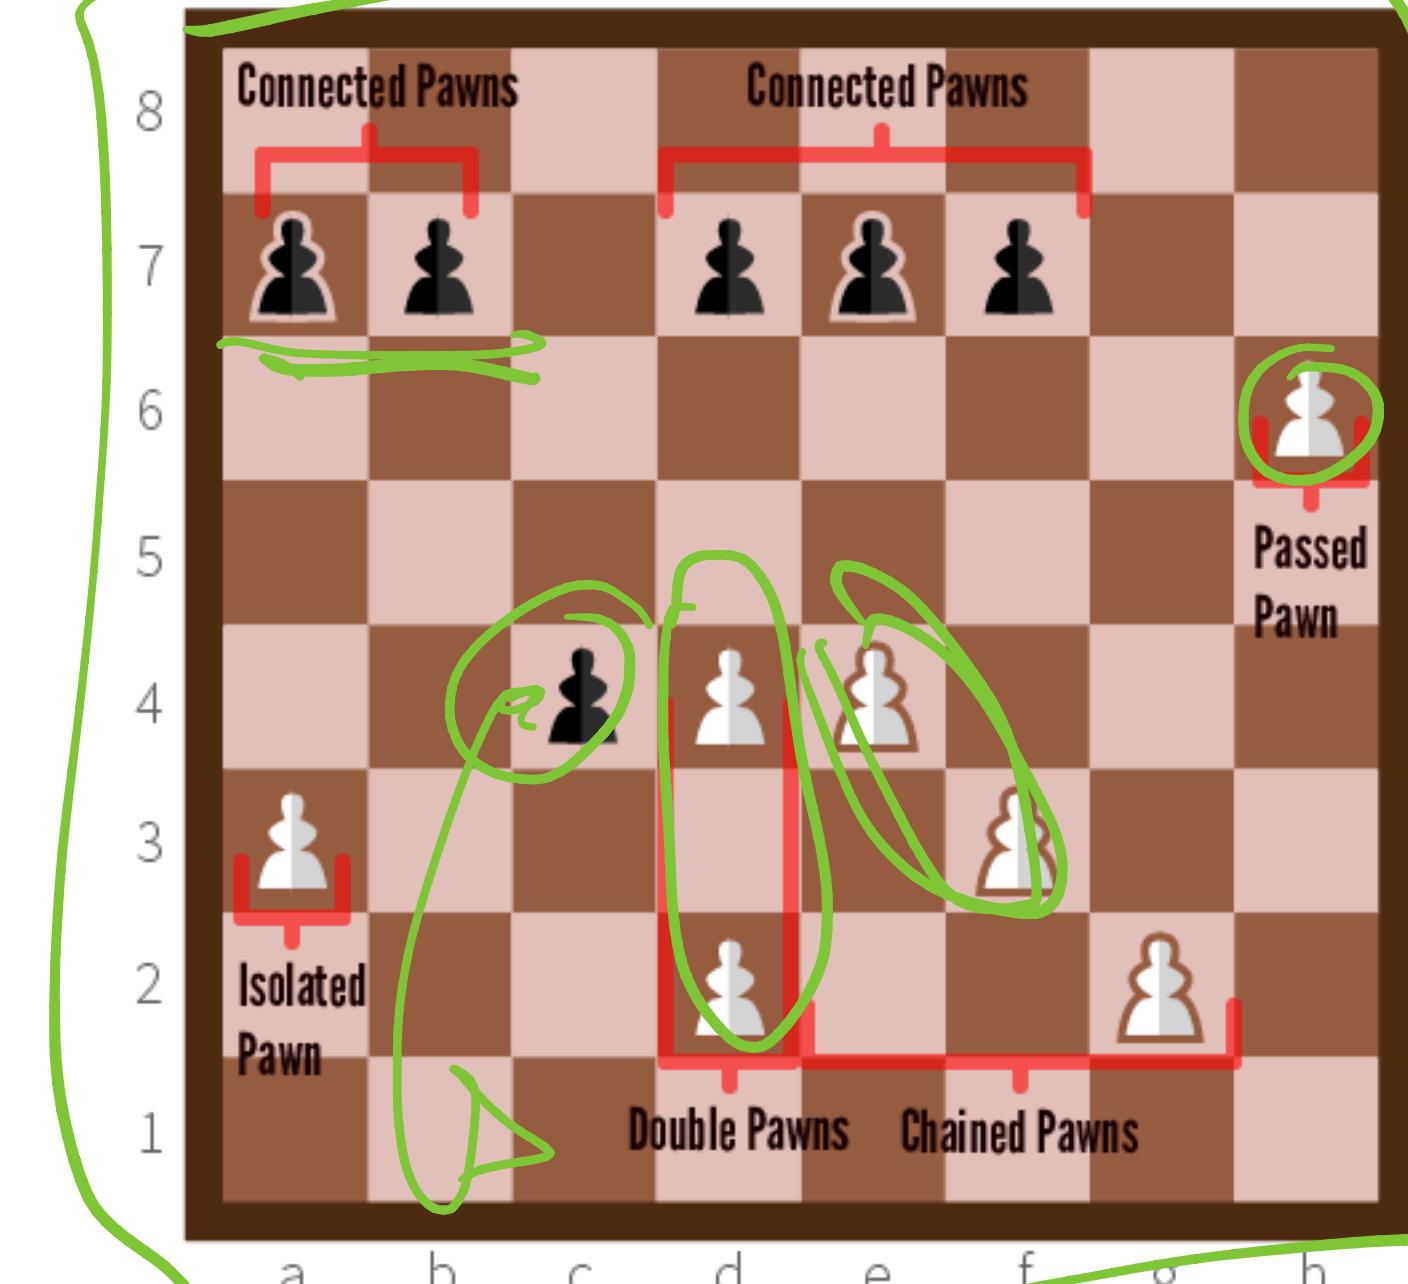
Possible Improvements

IBM
Beta Cut
Beta Pruning

- Heuristic Alpha-Beta Search

$$H\text{-MINIMAX}(s, d) = \begin{cases} EVAL(s) & \text{if } CUTOFF\text{-TEST}(s, d) = \text{true}, \\ \max_{a \in Actions(s)} H\text{-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if } \text{PLAYER}(s) = MAX, \\ \min_{a \in Actions(s)} H\text{-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if } \text{PLAYER}(s) = MIN. \end{cases}$$

Pieces and Point Value	
Pawn	1
Knight	3.
Bishop	3.
Rook	5
Queen	9
King	priceless



R

MinMax Algorithm

Possible Improvements

\leftarrow

- Use Iterative Deepening Depth-First Search
- Transposition Table 



R

MinMax Algorithm

Possible Improvements

- Monte Carlo Search

```
function MONTE-CARLO-TREE-SEARCH(state) returns an action
  tree  $\leftarrow$  NODE(state)
  while IS-TIME-REMAINING() do
    leaf  $\leftarrow$  SELECT(tree)
    child  $\leftarrow$  EXPAND(leaf)
    result  $\leftarrow$  SIMULATE(child)
    BACK-PROPAGATE(result, child)
  return the move in ACTIONS(state) whose node has highest number of playouts
```



↑

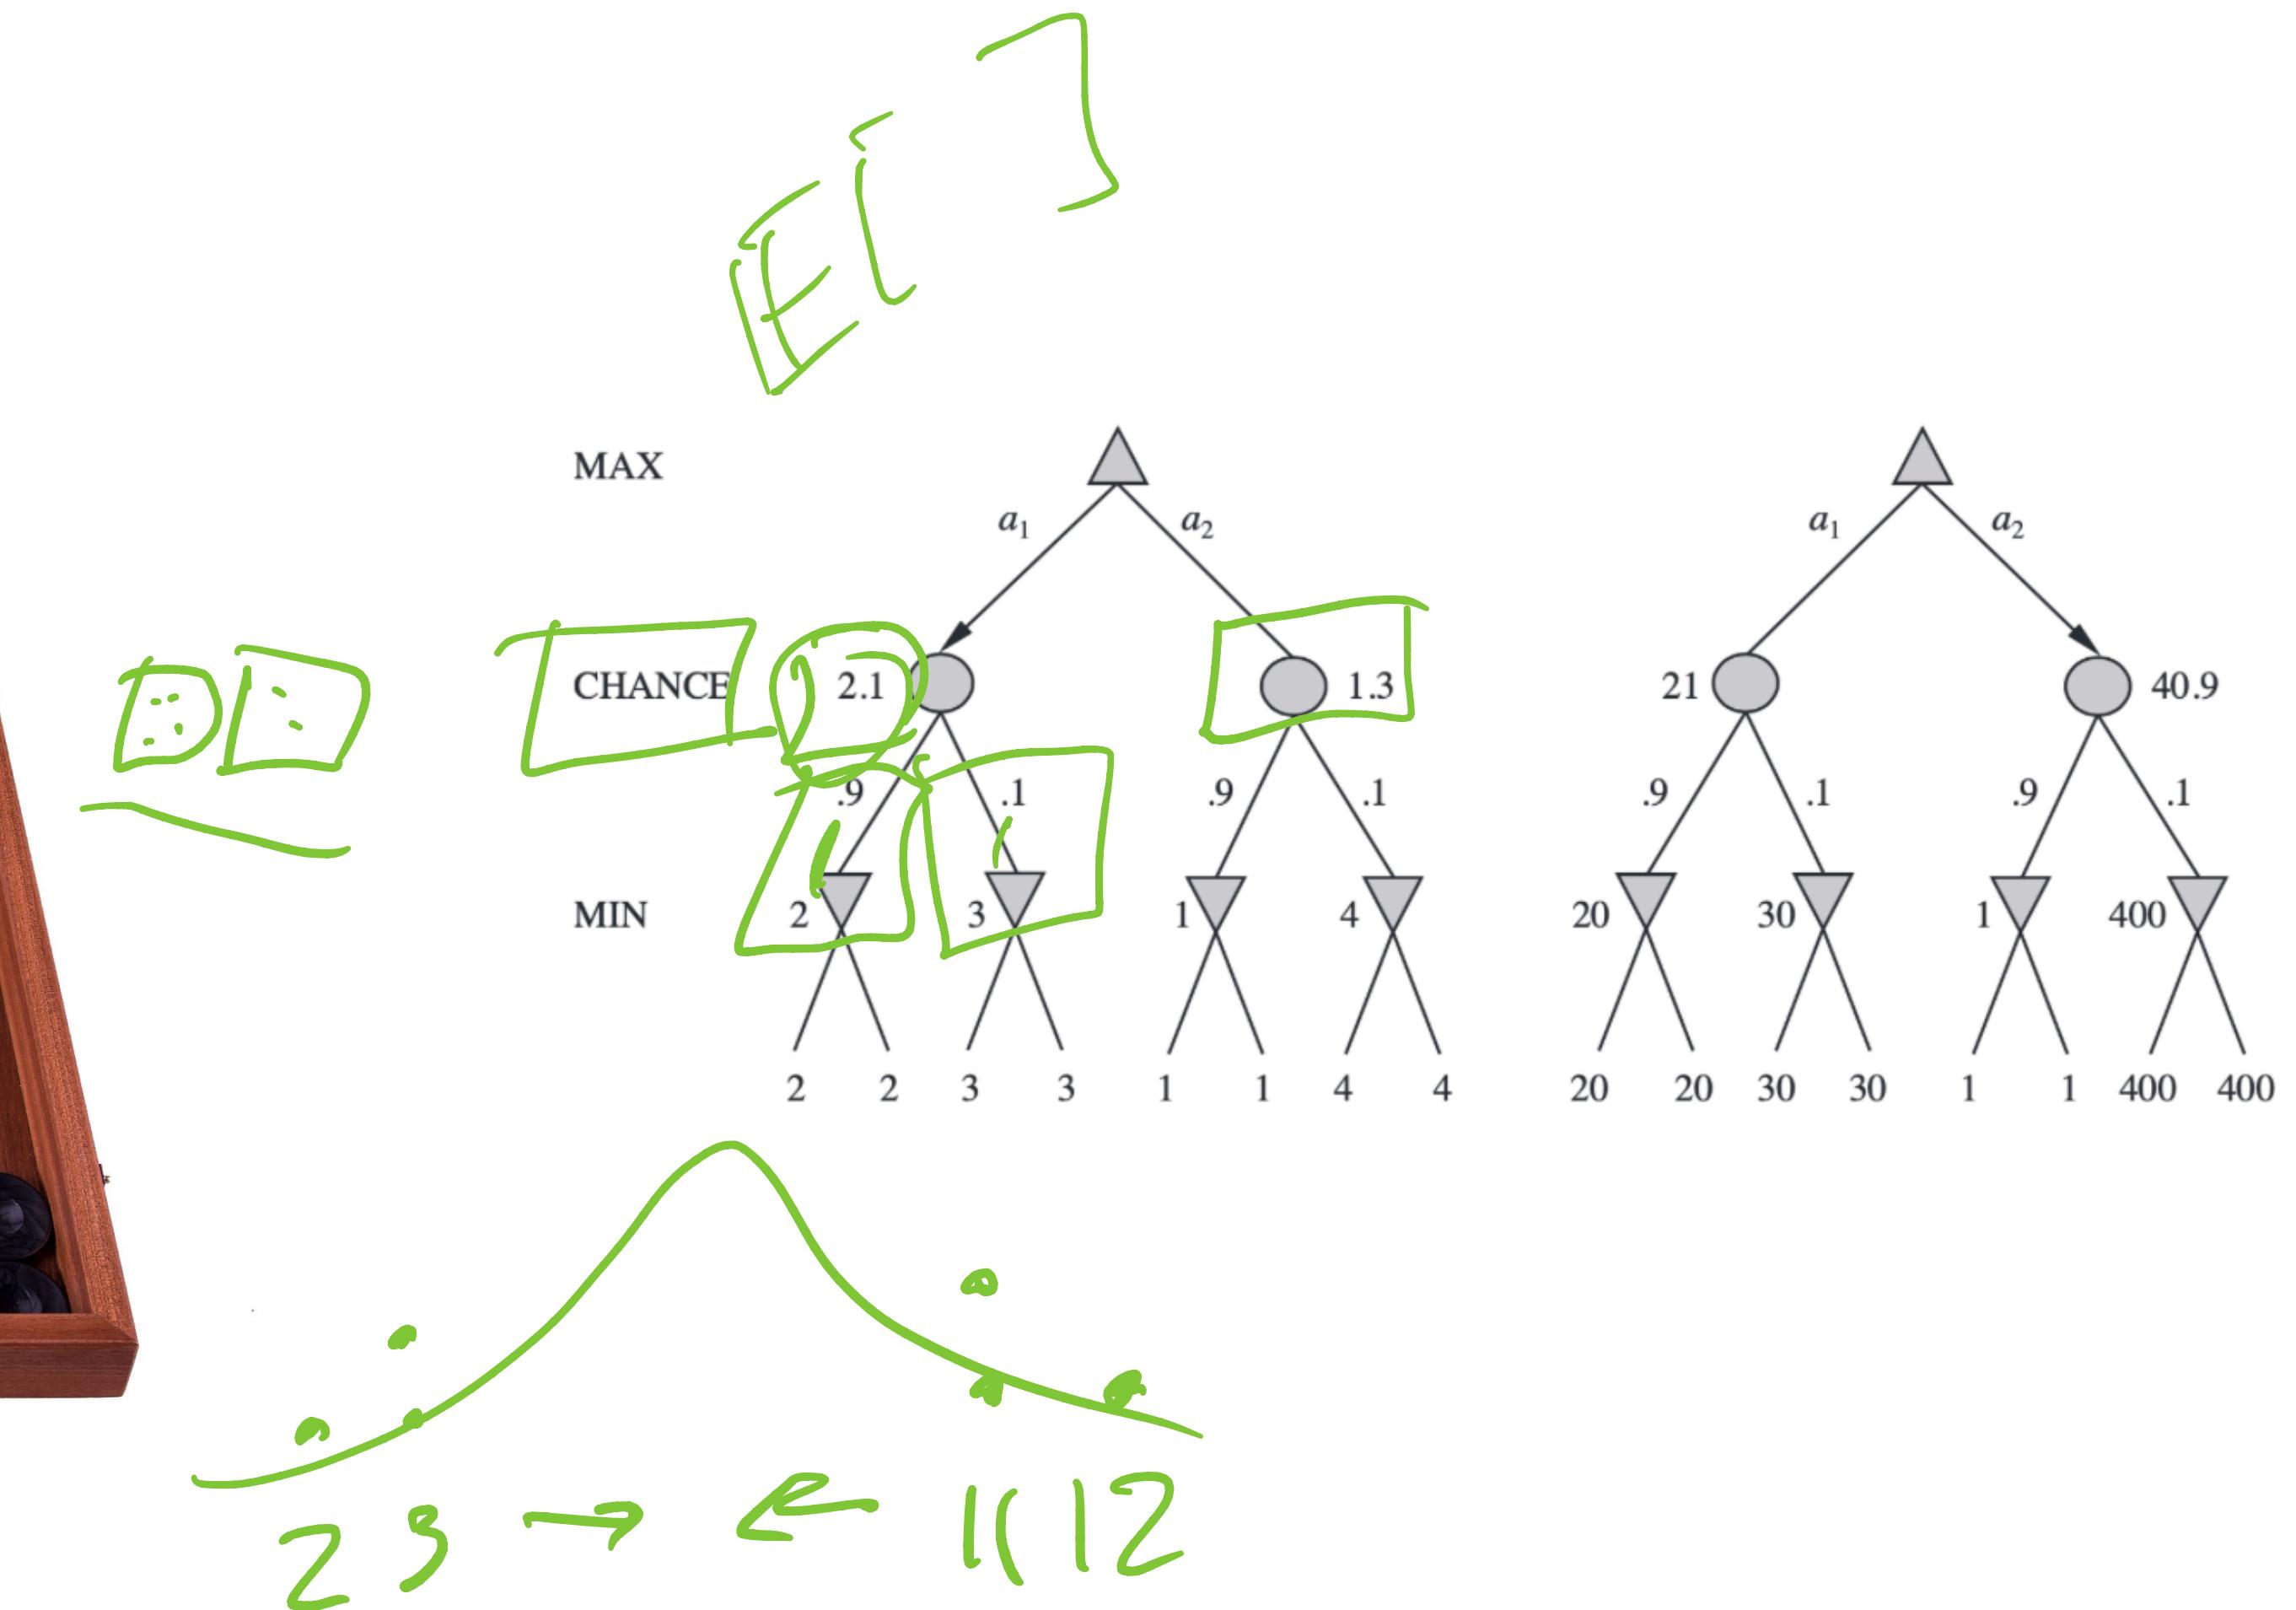
alpha
GO

R

MinMax Algorithm

Possible Improvements

- Stochastic Games



R

MinMax Algorithm Possible Improvements

• Stochastic Games



SHARE RESEARCH ARTICLE

[f](#) Superhuman AI for heads-up no-limit poker: Libratus beats top professionals

[t](#) [in](#) [g](#) [m](#) Noam Brown, Tuomas Sandholm *

* See all authors and affiliations

Science 26 Jan 2018:
Vol. 359, Issue 6374, pp. 418-424
DOI: 10.1126/science.aao1733

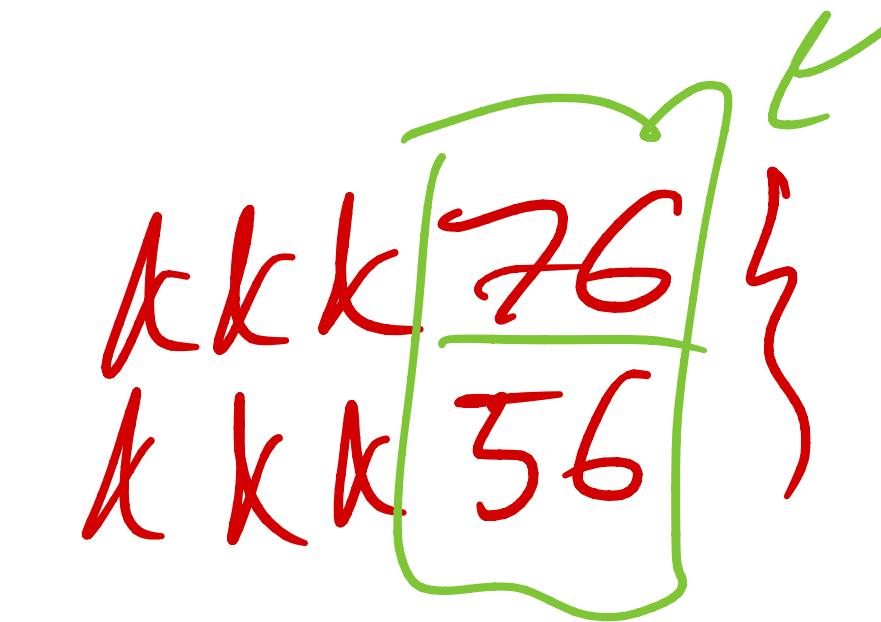
Article Figures & Data Info & Metrics eLetters PDF

Figures ▾Tables ▾Supplementary Materials ▾Additional Files

Fig. 1 Subgame solving.

(Top) A subgame (red) is reached during play. Blue and red indicate the blueprint strategy. The white path indicates the action sequence before the reached subgame. (Middle) A more-detailed strategy for that subgame is determined by solving an augmented subgame in which, on each iteration, the opponent is dealt a random poker hand and given the choice of taking the expected value of the old abstraction (red) or of playing in the new, finer-grained abstraction (green), where the strategy for both players can change. This forces Libratus to make the finer-grained strategy at least as good as that in the original abstraction against every opponent poker hand. (Bottom) The new strategy is substituted in place of the old one.

[Download high-res image](#) [Open in new tab](#) [Download Powerpoint](#)

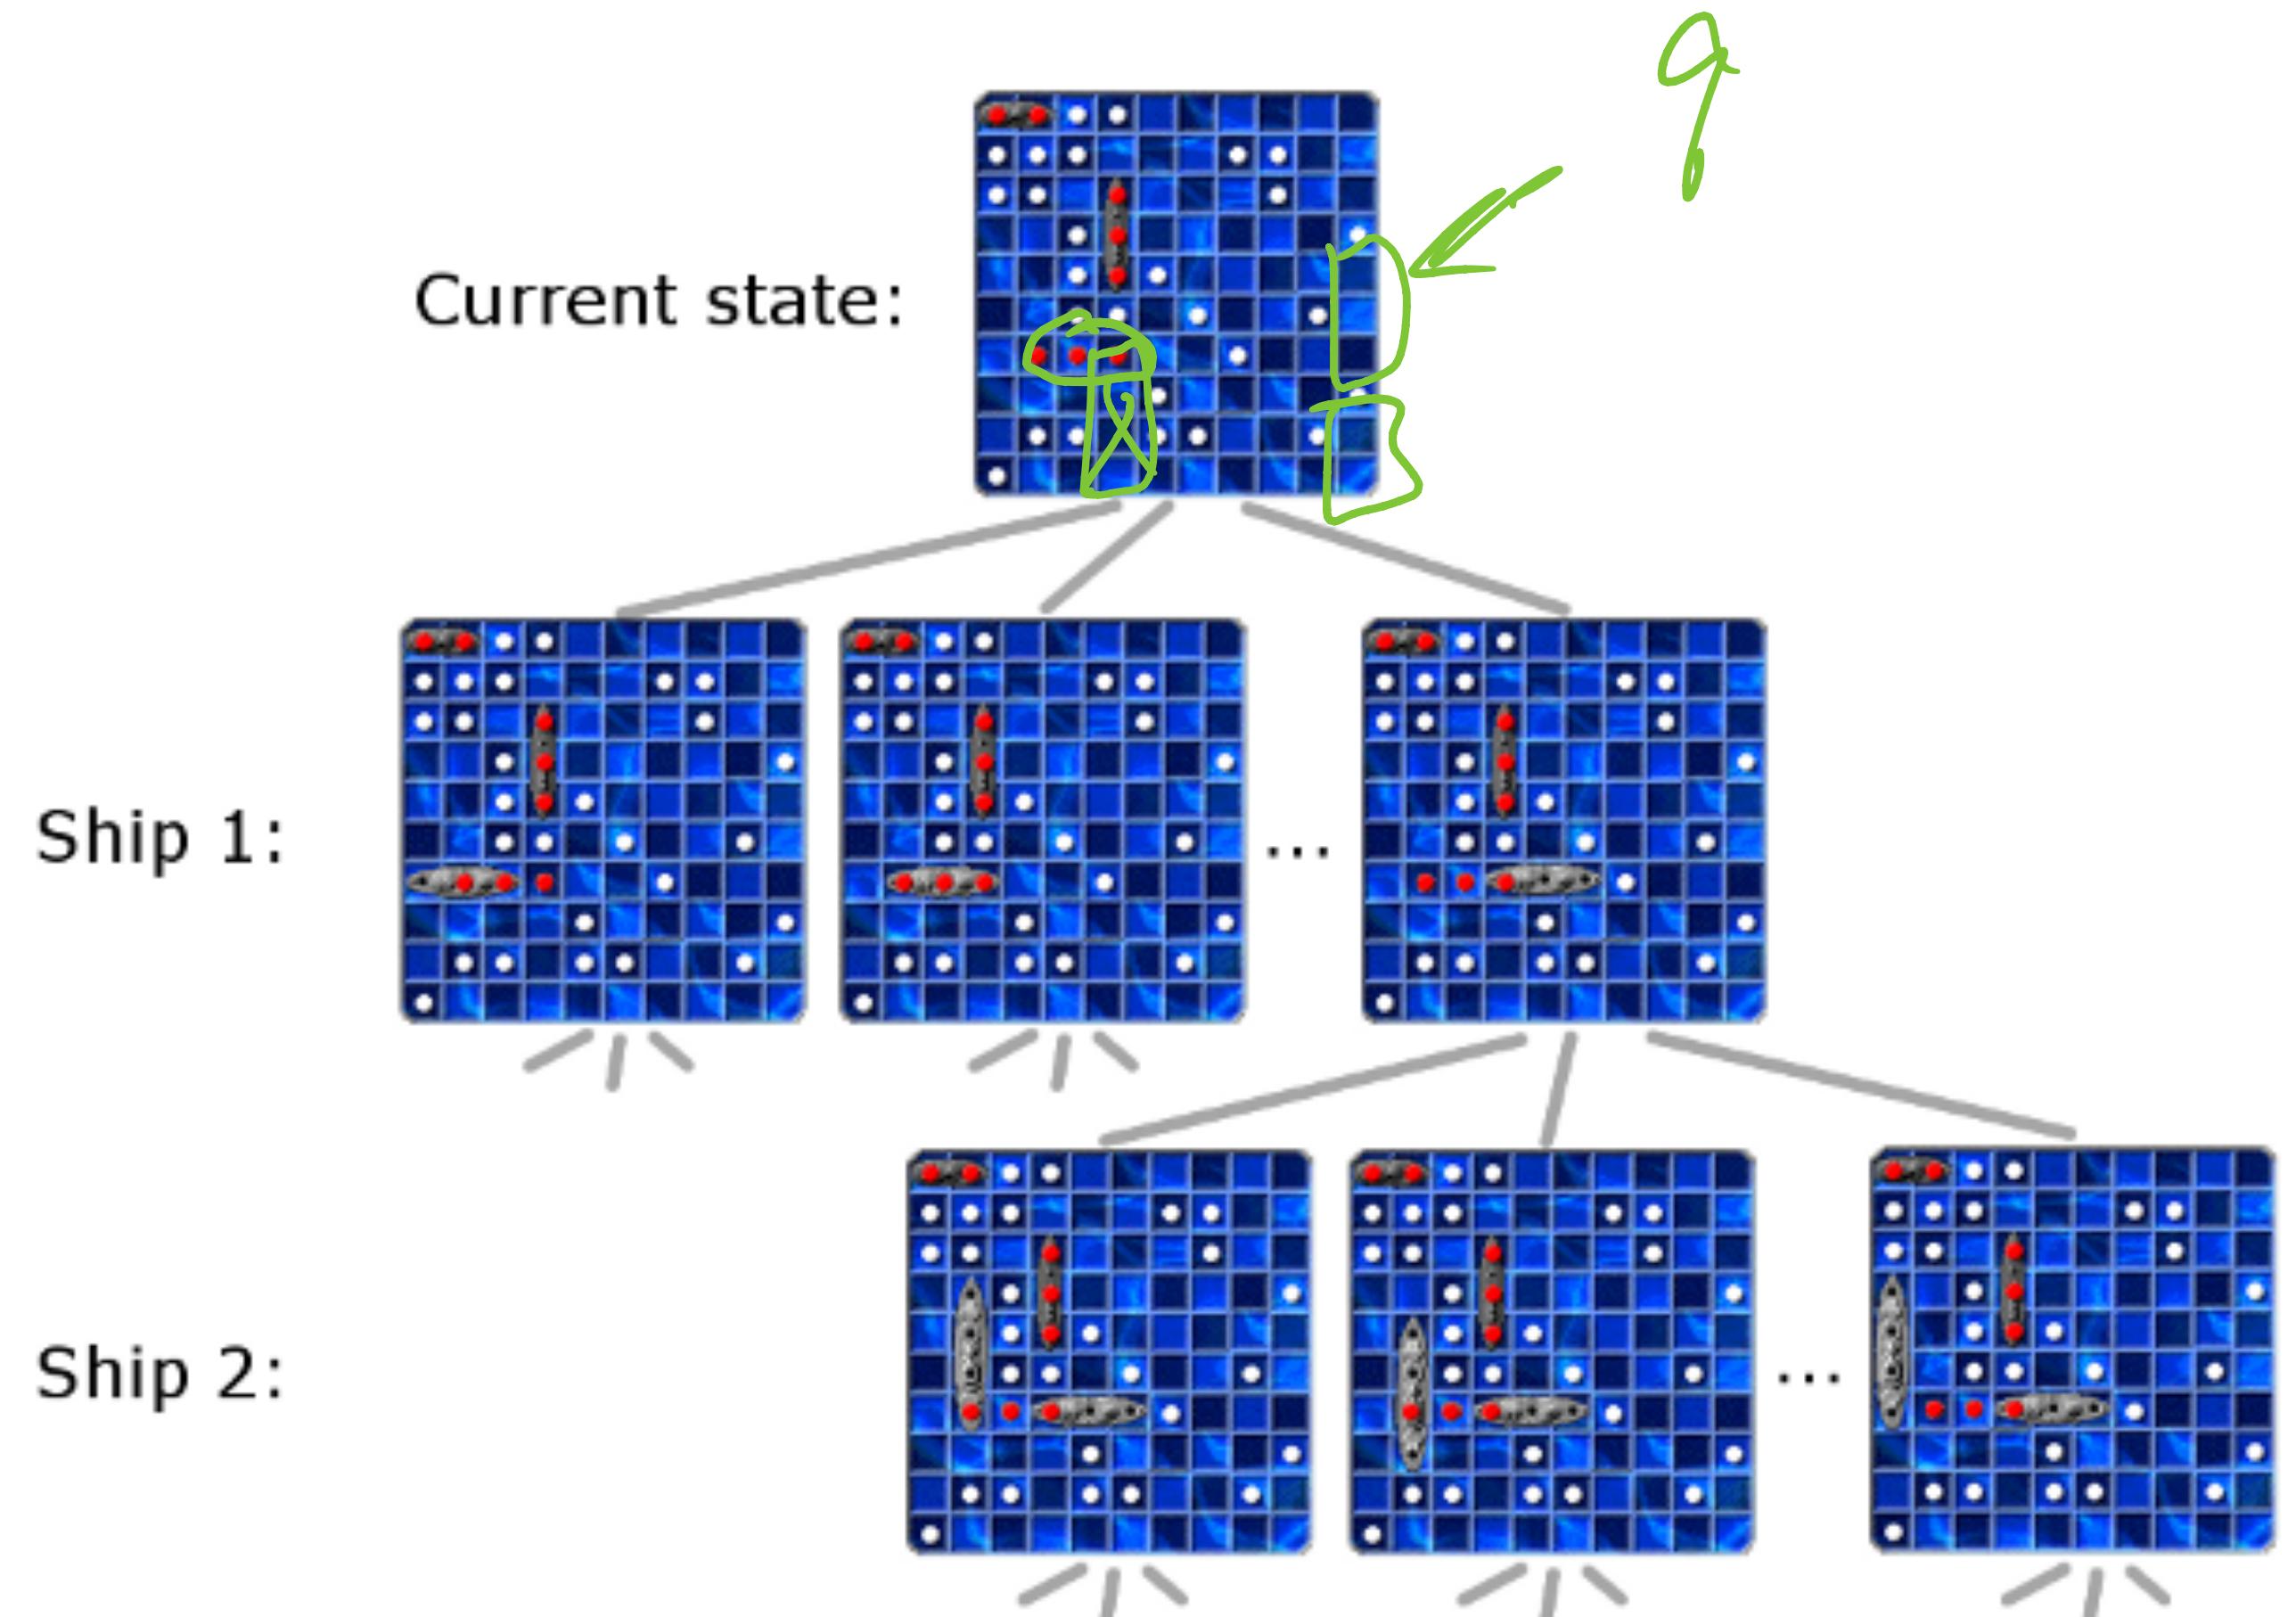


R

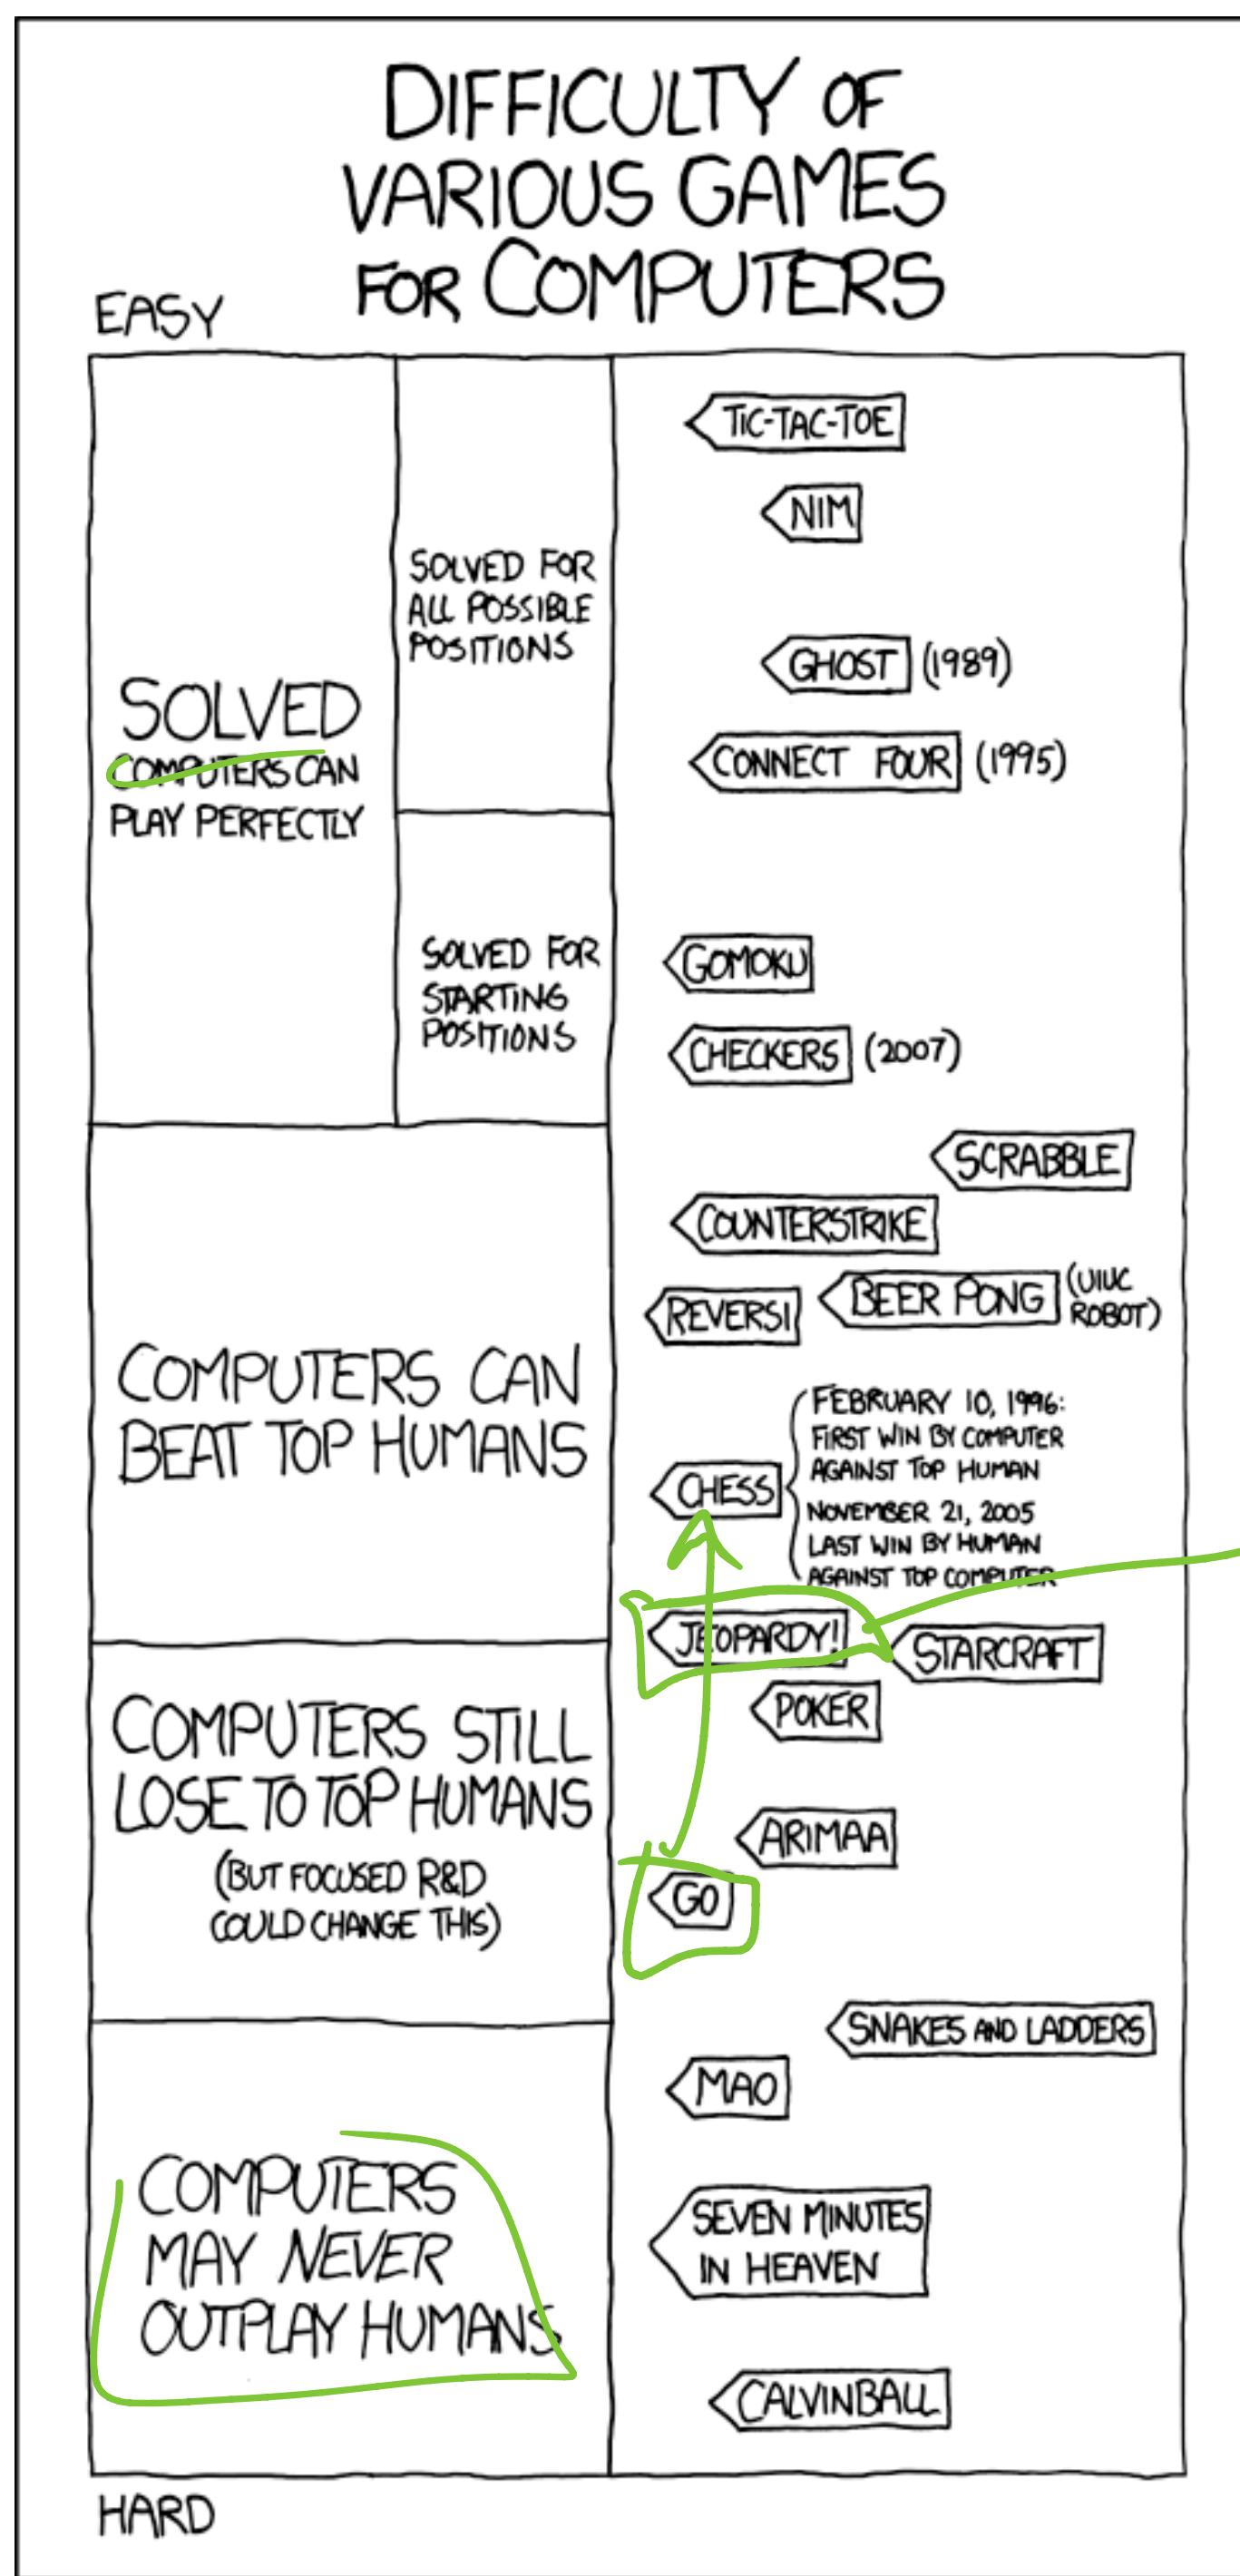
MinMax Algorithm

Possible Improvements

- Partial Observability



R



State of the art

IBM
Watson